

Version 1.07

UCSD CSE 30

Computer Organization and Systems Programming

Aarch32 Assembly - Introduction

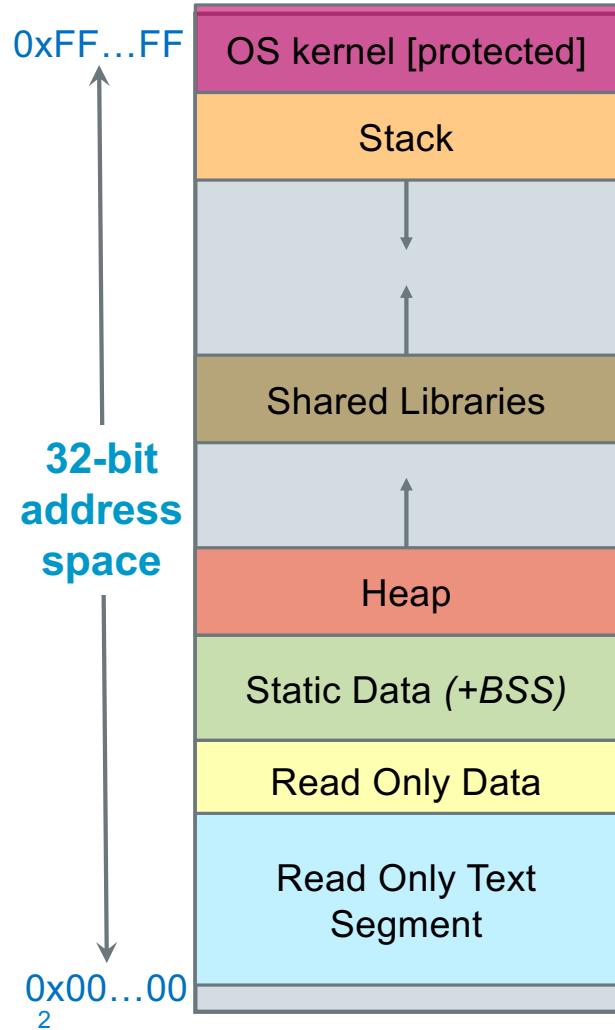
Week 6

Lecture 16

Keith Muller



# Assembly and Machine Code



- Machine Language (or code): Set of instructions the CPU executes are encoded in memory using patterns of ones and zeros
- Assembly language is a symbolic version of the machine language
- Each assembly statement (called an **Instruction**), executes exactly **one** from a list of simple commands
  - Instructions describe operations (e.g., =, +, -, \*)
- Each line of arm32 assembly code contains at most one instruction
- Assembler (gnu as) translates assembly to machine code

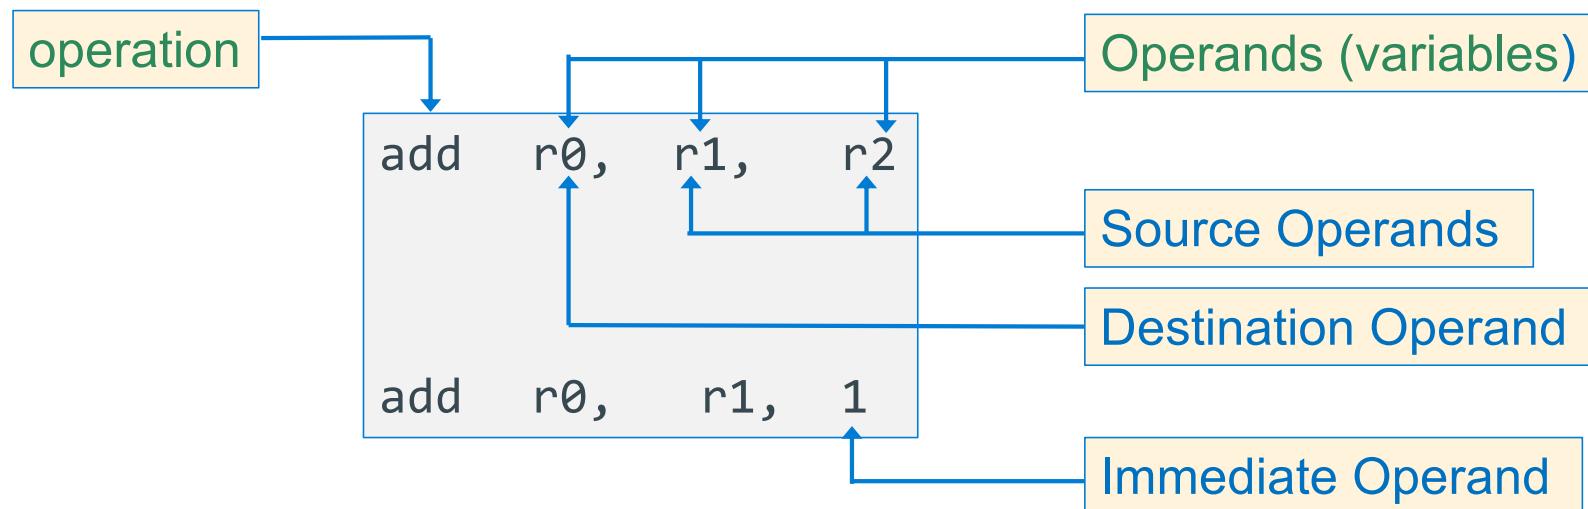
Memory Address	word (4-bytes) contents	Assembly Language
1040c:	e28db004	add fp, sp, 4
10410:	e59f0010	ldr r0, [pc, 16]
10414:	ebfffffb3	bl 102e8 <printf>
10418:	e3a00000	mov r0, 0
1041c:	e24bd004	sub sp, fp, 4

high <- low bytes

Machine Code

# Anatomy of an Assembly instruction

- Assembly language instructions specify an **operation** and the **operands** to the instruction (arguments of the operation)
- Three basic types of **operands**
  - **Destination**: where the result will be stored
  - **Source**: where data is read from
  - **Immediate**: an actual value like the **1** in  $y = x + 1$



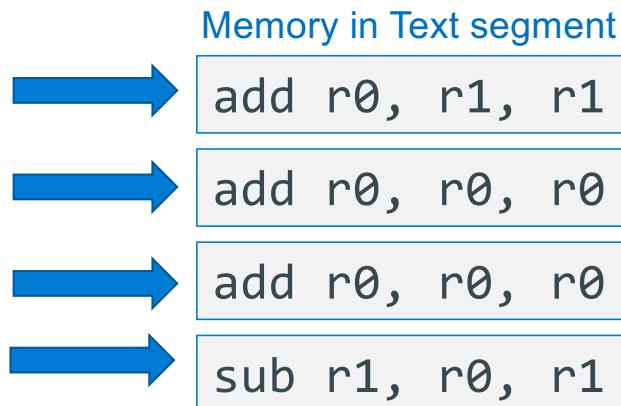
# Meaning of an Instruction

- Operations are abbreviated int **opcodes** (1– 5 letters)
- **Assembly Instructions** are specified with a very regular syntax
  - Opcodes are followed by **arguments**
  - Usually the **destination argument is next**, then **one or more source arguments** (this is not strictly the case, but it is generally true)
- Why this order?
- Analogy to C or Java

```
int r0, r1, r2;  
r0 = r1 + r2;  
  
/* r0 = r1 + r2 */  
add    r0,   r1,   r2
```

# Program Execution: A Series of Instructions

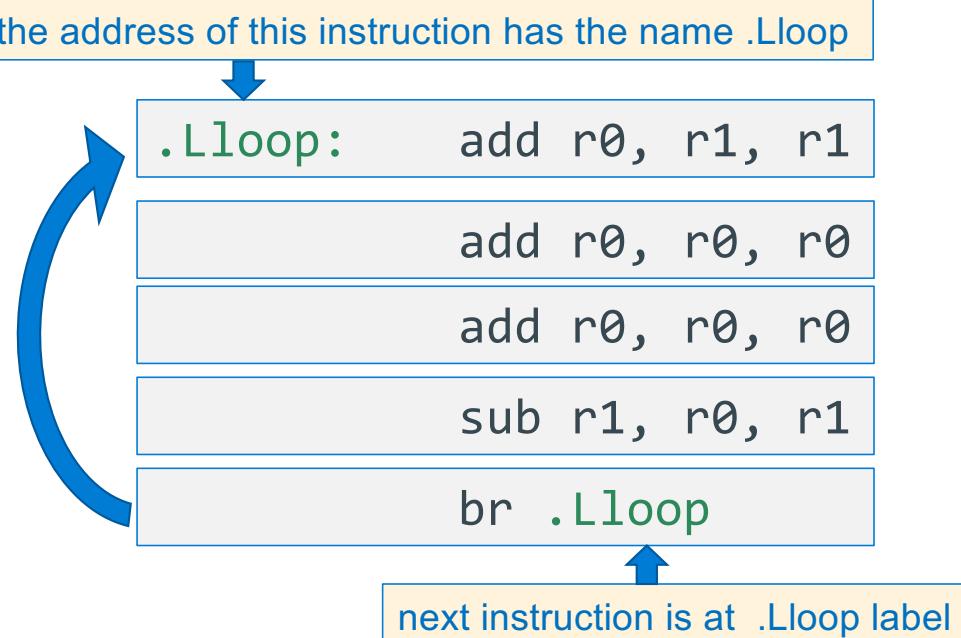
- Generally (there are exceptions like function calls, loops etc.)
- Instructions are retrieved sequentially from memory
- An instruction executes to completion before the next instruction is completed
- But there are exceptions to this (later and in cse 141, cse 142)



r0 = 1 r1 = 2
r0 = 4 r1 = 2
r0 = 8 r1 = 2
r0 = 16 r1 = 2
r0 = 16 r1 = 14

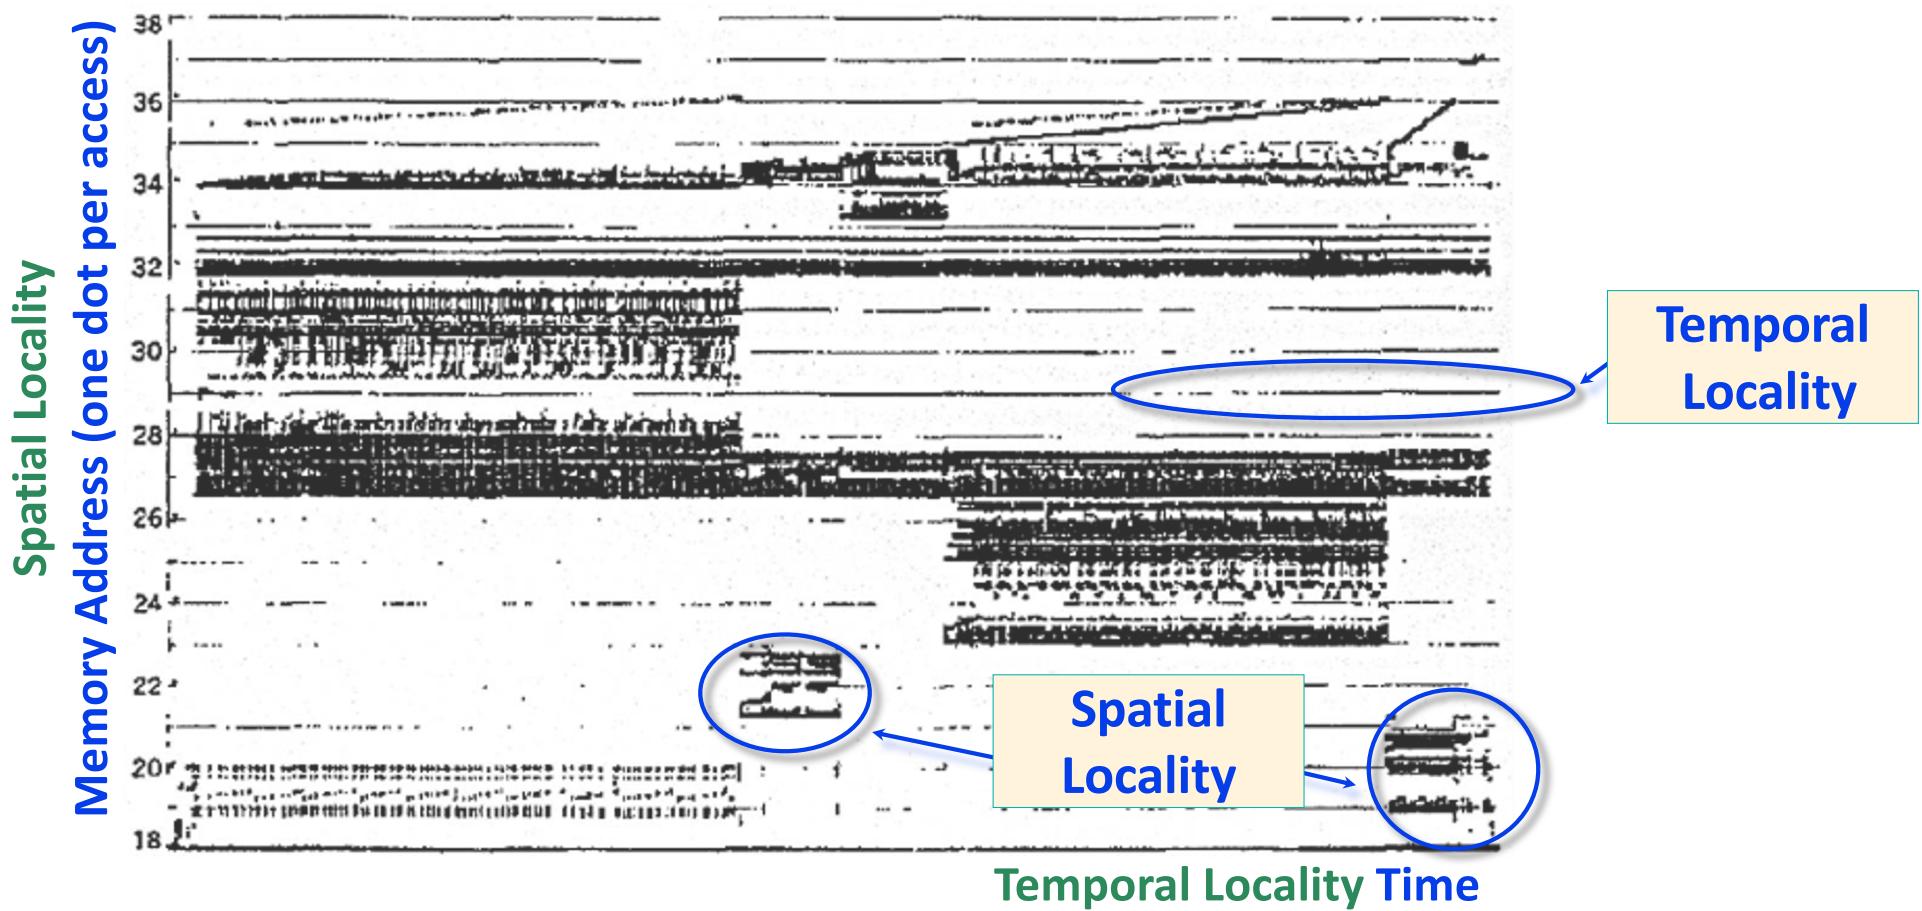
# Program Execution: Looping in the Execution Flow

- Repeat the series of instructions in a loop means **altering the flow of execution**
- This is used with if statements and loops
- Below is an example of an **infinite** loop



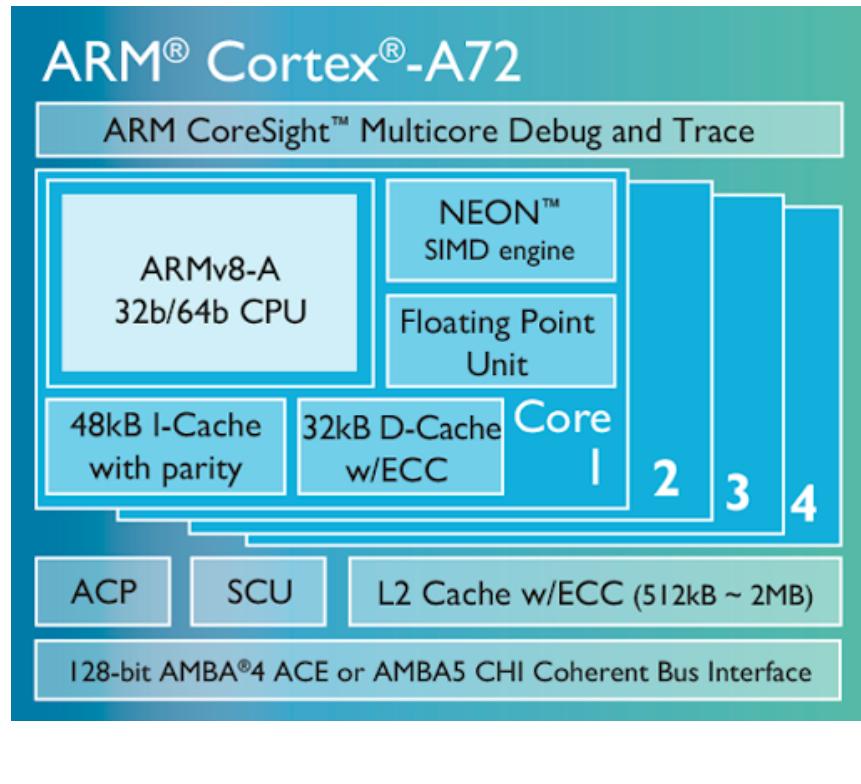
$r0 = 1 \ r1 = 2$
$r0 = 4 \ r1 = 2$
$r0 = 8 \ r1 = 2$
$r0 = 16 \ r1 = 2$
$r0 = 16 \ r1 = 14$

## Principle of Locality: Spatial and Temporal Memory Access of Executing Programs

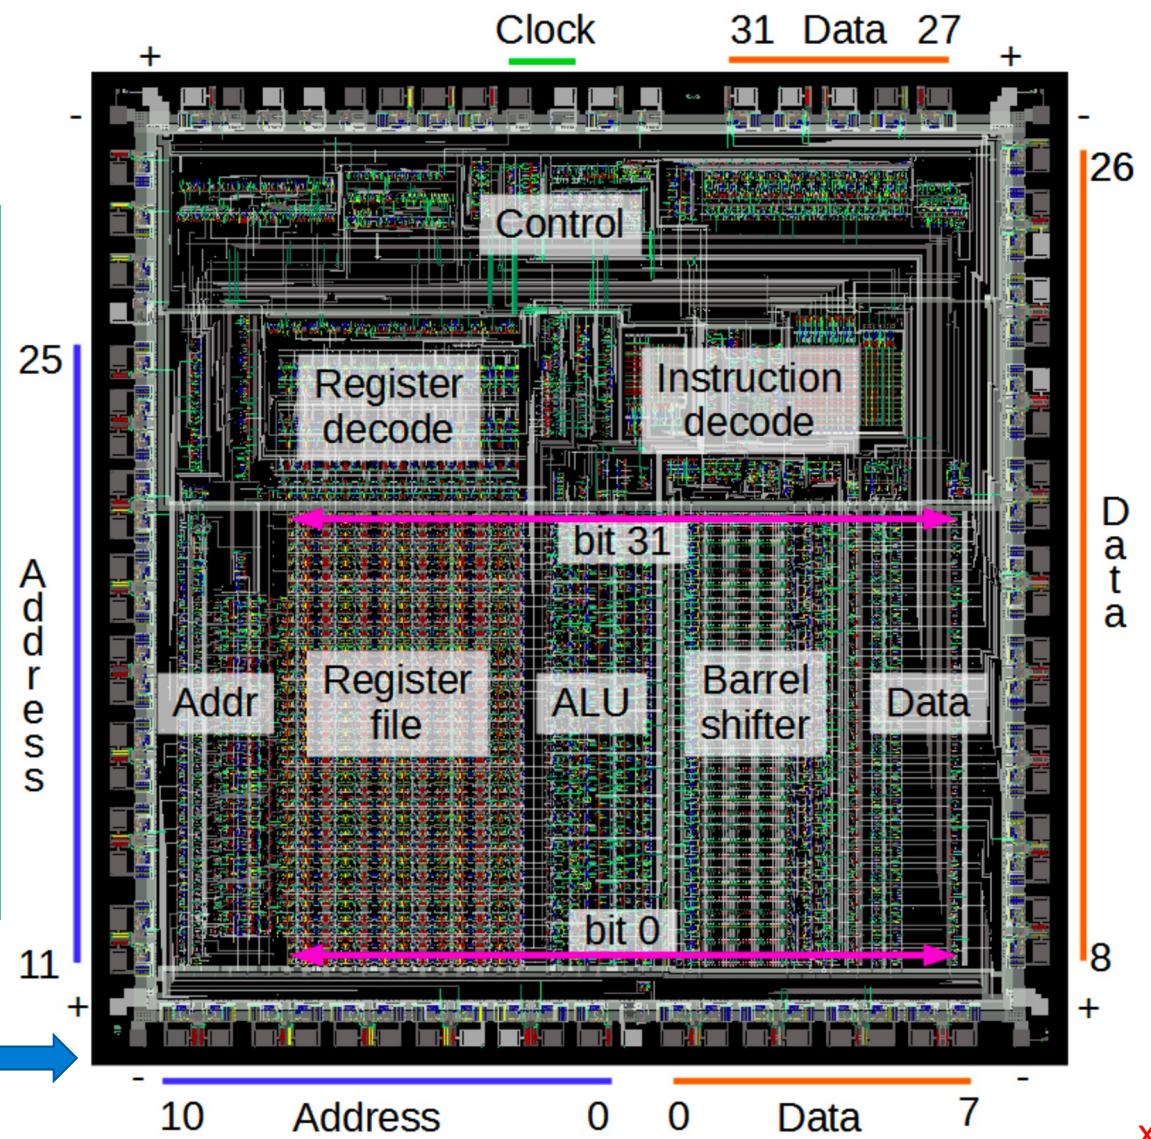


7 Donald J. Hatfield, Jeanette Gerald: Program Restructuring for Virtual Memory. IBM Systems Journal 10(3): 168-192 (1971) X

# Arm Core Organization & Floorplan Examples



Single core **arm** die (not an A72) **Floorplan**



Version 1.07

# UCSD CSE 30

## Computer Organization and Systems Programming

### Aarch32 Assembly - Introduction

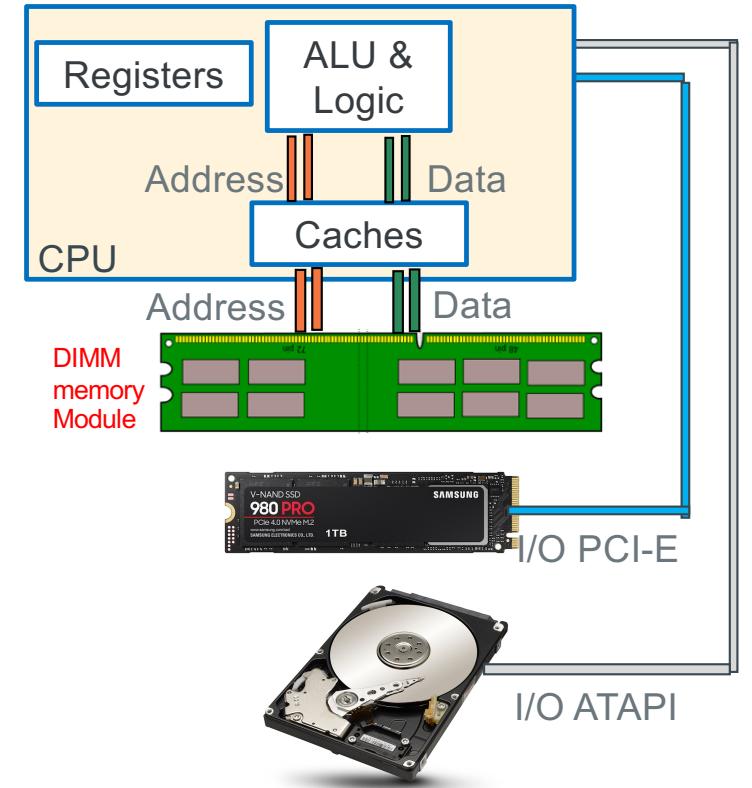
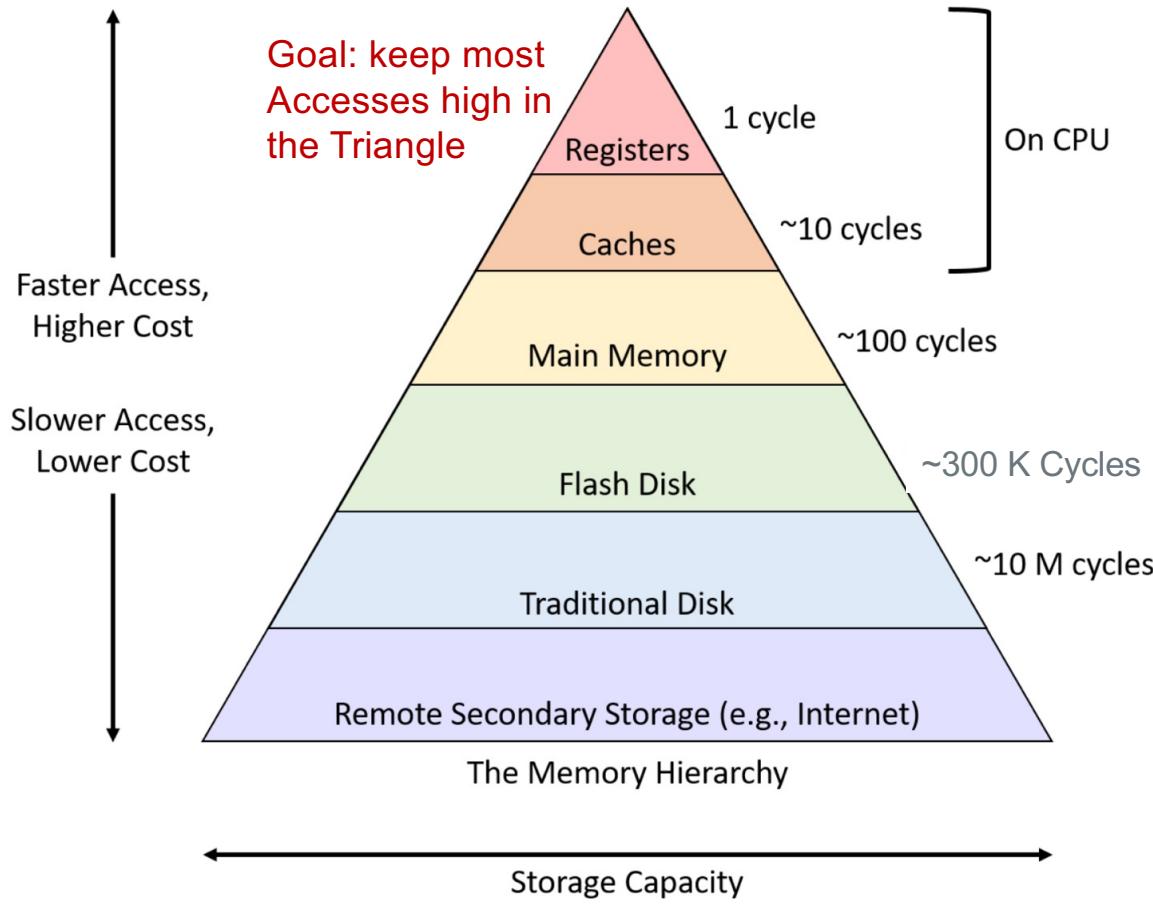
Week 6

Lecture 17

Keith Muller

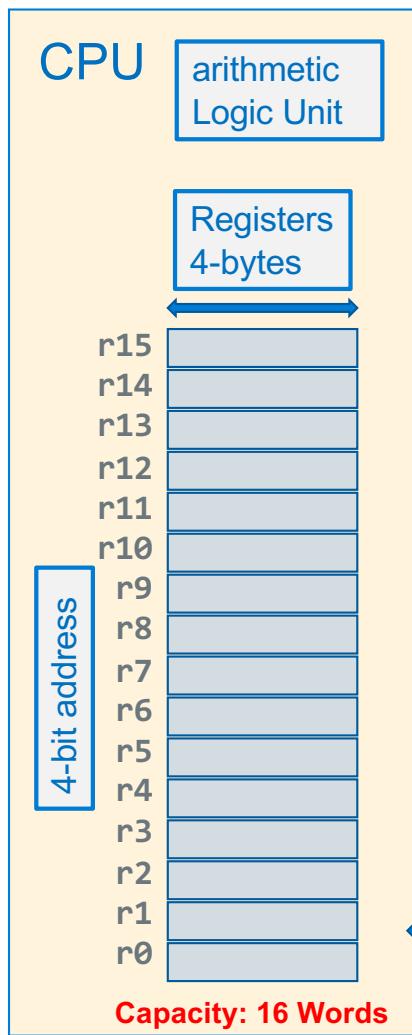


# Memory Triangle: Taking Advantage of Locality for Cost/Performance



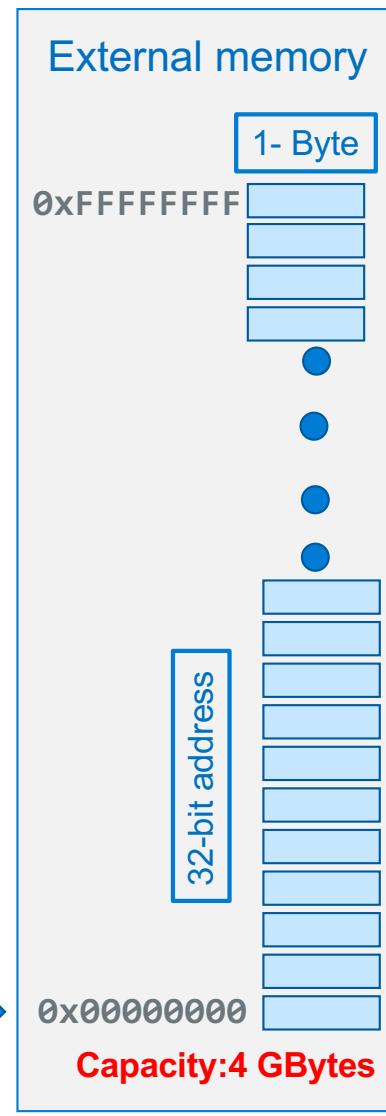
Assume 1 clock cycle  
per machine instruction

## 32-Bit Arm - Registers



- Registers are **memory located within the CPU**
- Registers are the **fastest** read and write storage
- Register is **word size in length** stores 32-bit values
  - Memory is accessed using **pointers** in registers
- In assembly language the registers have **predefined names** starting with an **r** to differentiate them from **memory addresses** which are **labels** (**address**)
- 16 registers: from **r0** to **r15** (encoded: **0x0 – 0xf**)

CPU Memory Bus = Address + Data

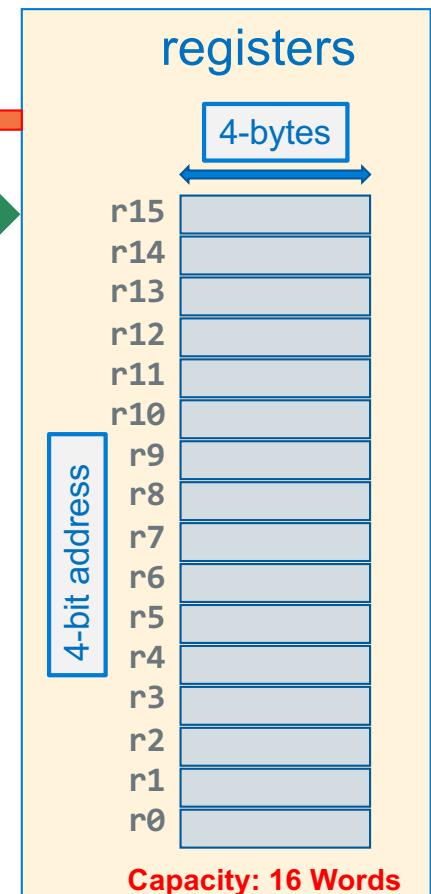
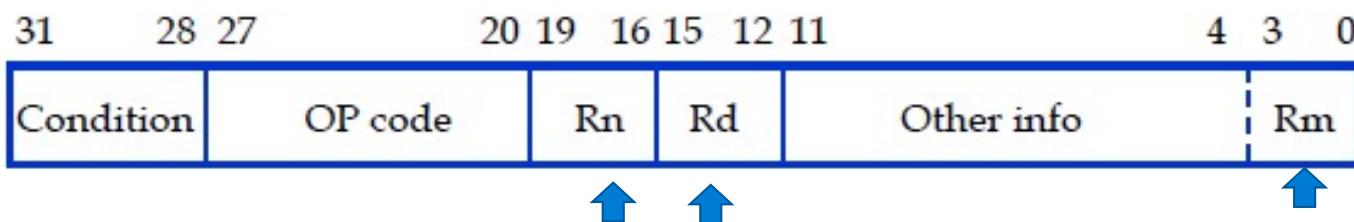


# 32-Bit Arm - Registers

All computations (add, subtract, etc.) are performed in the ALU

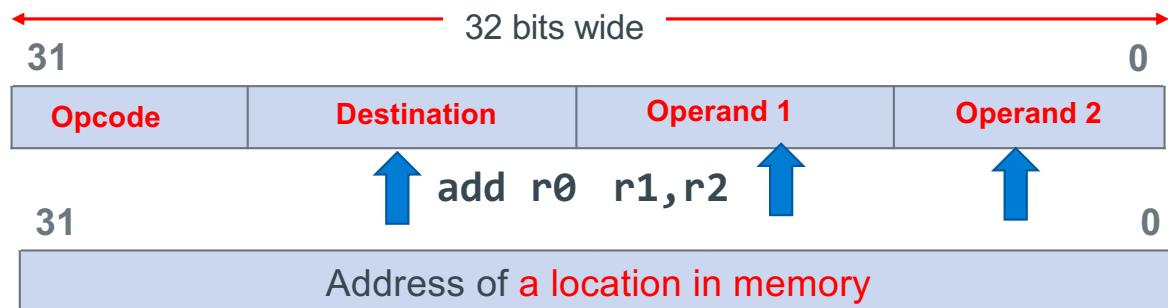
Arithmetic & Logic Unit (ALU)

- Almost all arithmetic, logic operations and data movement operations involve at least one register
- As a result, Register addresses are **directly encoded** into 4-bit fields in machine instructions (see below)

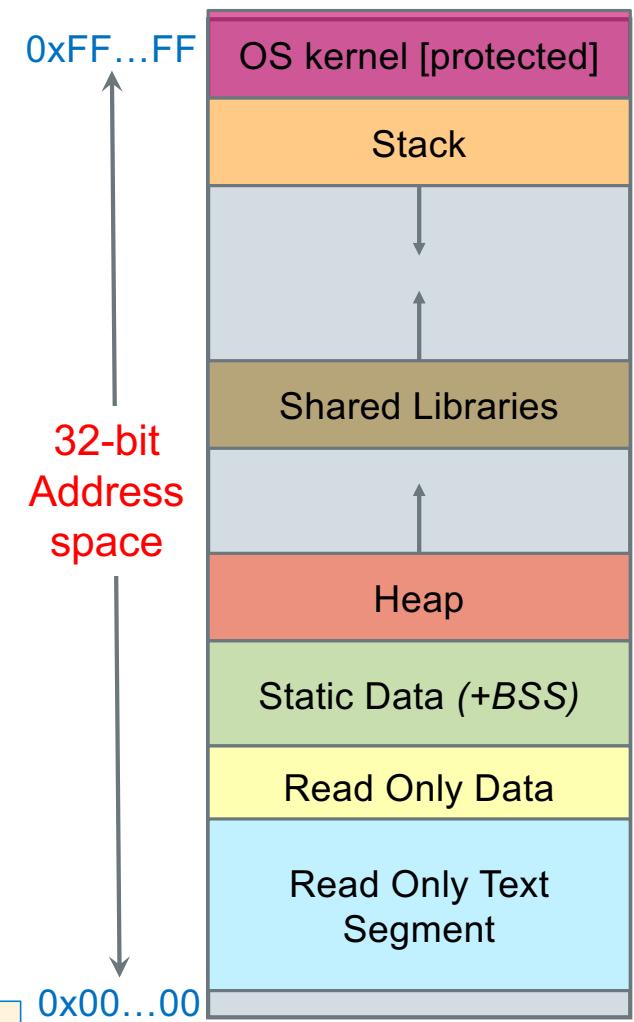


## How to Access Memory?

- Consider  $r0 = r1 + r2$ 
  - Operation code: add
  - Operand 1:  $r1$
- Aarch32 Instructions are always word size: 32 bits wide
  - Some bits must be used to specify the operation code
  - Some bits must be used to specify the destination
  - Some bits must be used to specify the operands
- Address space is 32 bits wide – **POINTERS** in registers



**NOT ENOUGH BITS for FULL Addresses in the instruction**



## Using Registers as Pointers to Memory - Load

We want to do a  $x[1] = x[0]$

We have to do this in two steps

```
int r3;  
int x[2];  
int *r1 = &x; // r1 contains address  
...  
r3 = *(r1); // memory to register  
*(r1 + 1) = r3; // register to memory
```

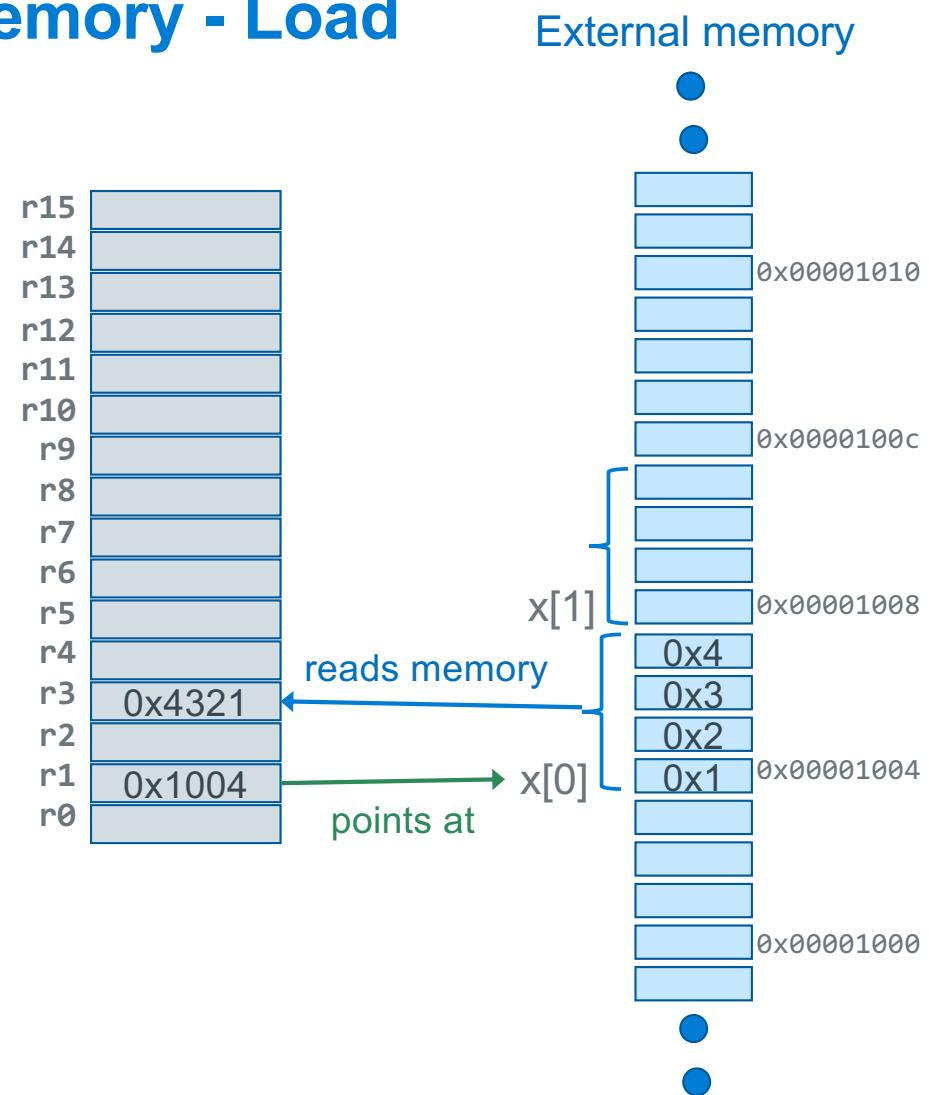
Load register from memory

```
ldr r3, [r1, 0]
```

address =  $r1 + 0 = 0x1004$

The [] around the operands is like the  
\* dereference op

we will cover this instruction in more detail later



# Using Registers as Pointers to Memory - Store

We want to do a  $x[1] = x[0]$

We have to do this in two steps

```
int r3;  
int x[2];  
int *r1 = &x; // r1 contains address  
...  
r3 = *(r1); // memory to register  
*(r1 + 1) = r3; // register to memory
```

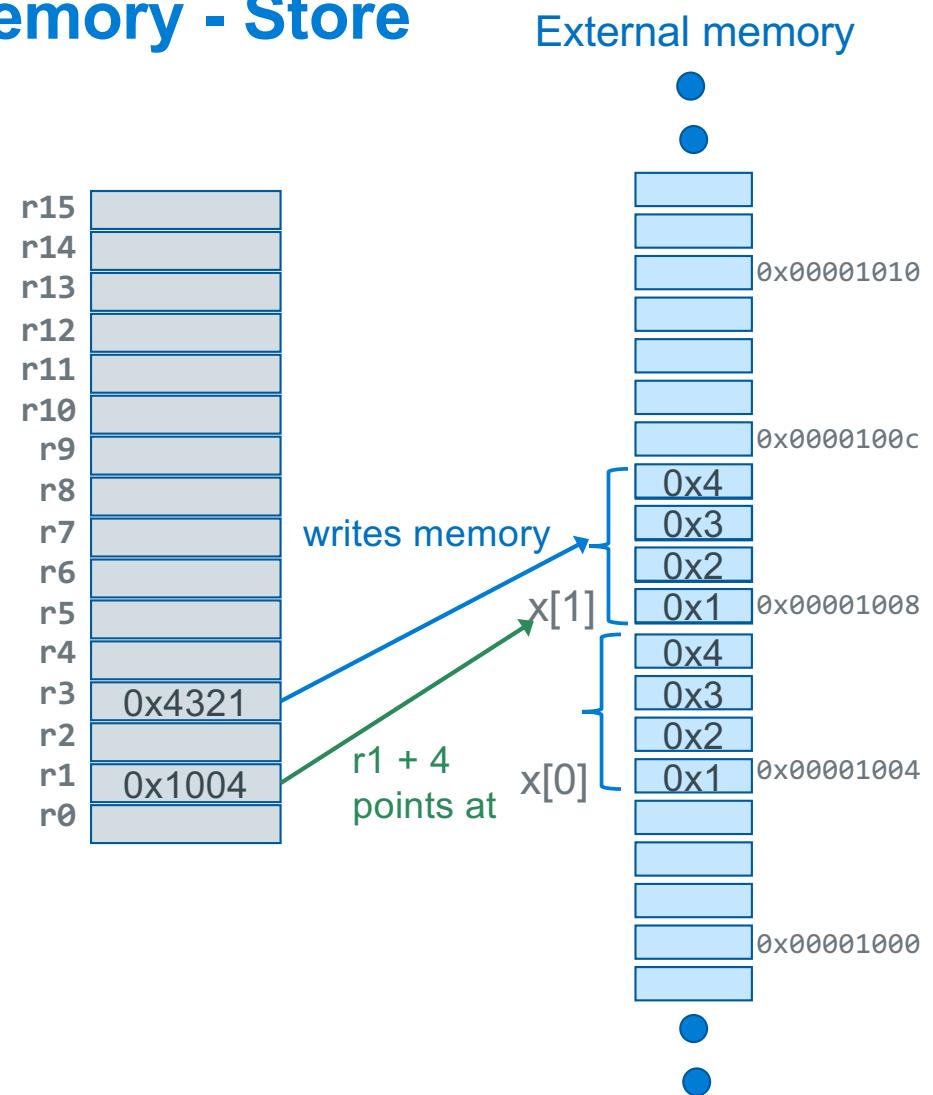
Store register to memory

str r3, [r1, 4]

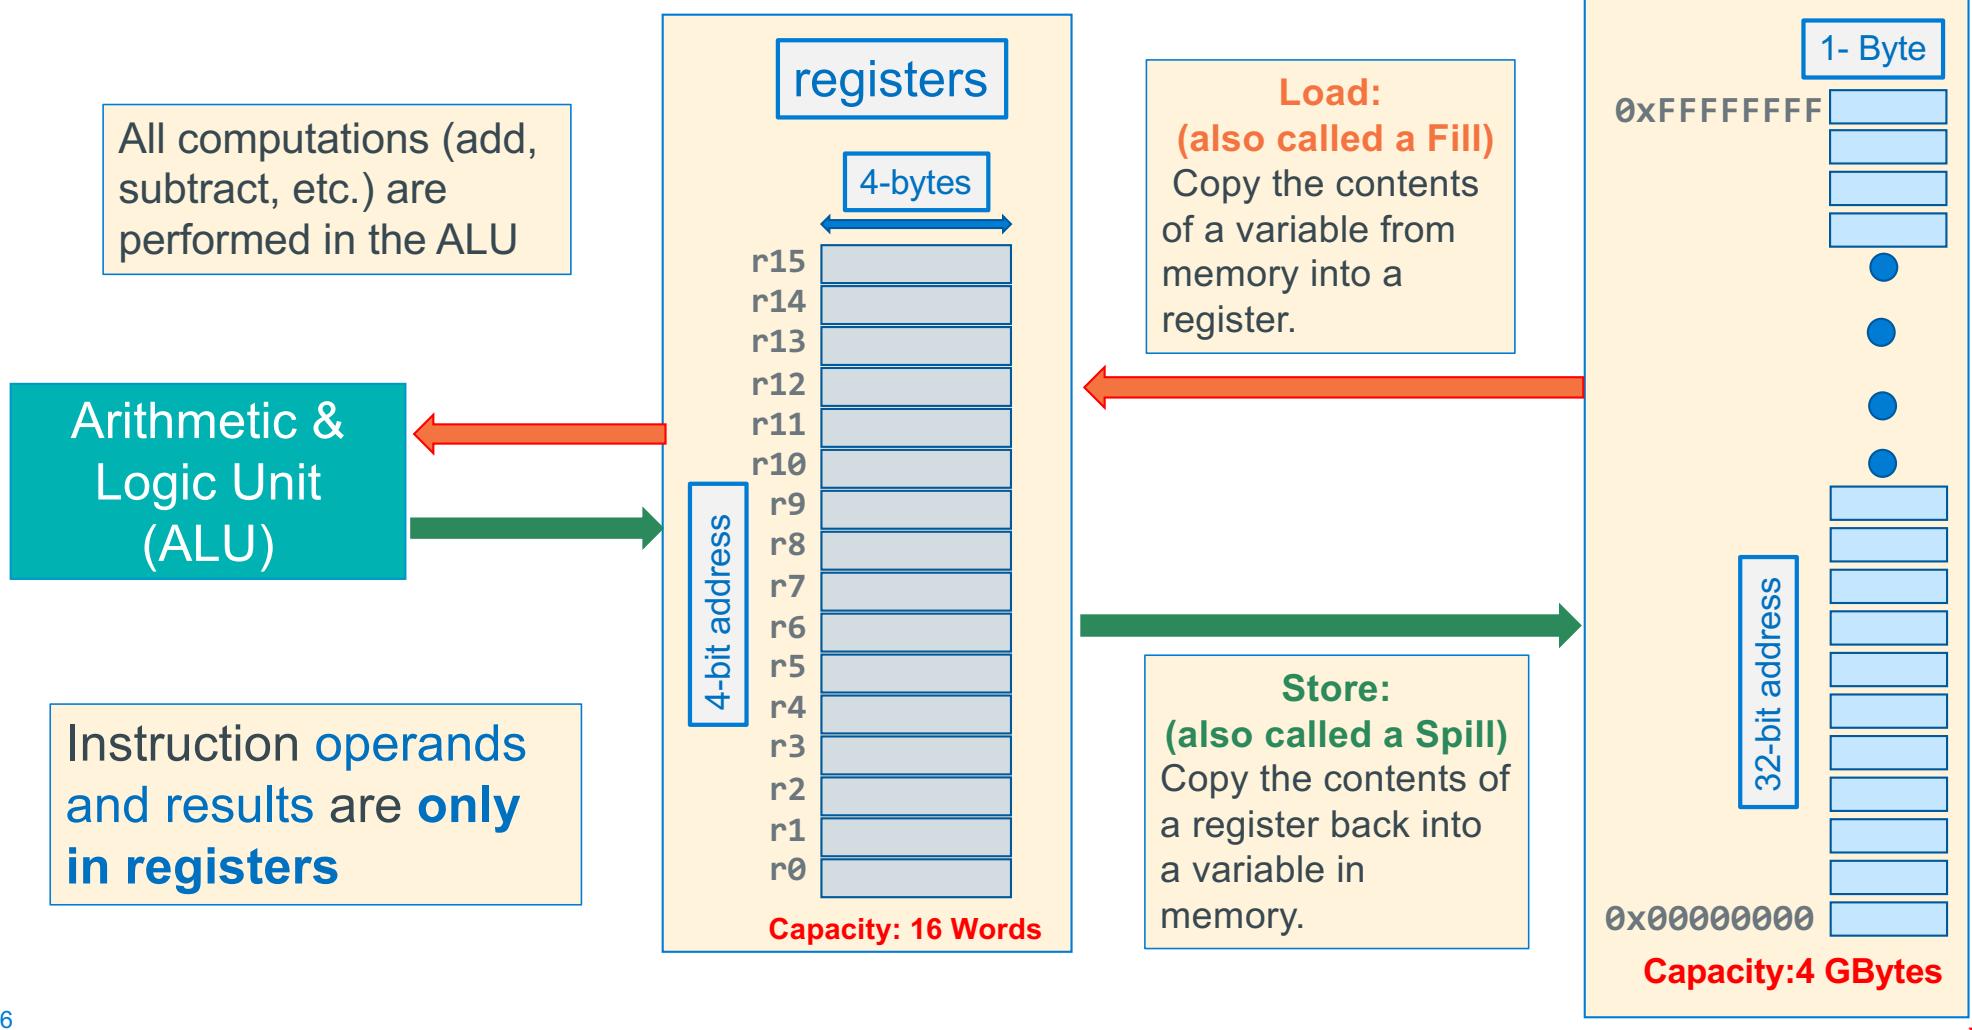
address =  $r1 + 4 = 0x1008$

The [ ] around the operands is like the  
\* dereference op

We will cover this instruction in more detail later



# 32-Bit Arm is a Load/Store Architecture



## Arm Register Summary

- 16 Named registers r0 – r15
- The operands of almost all instructions are registers
- To operate on a variable in memory do the following:
  1. Load the value(s) from memory into a register
  2. Execute the instruction
  3. Store the result back into memory (only if needed!)
- Going to/from memory is expensive
  - 4X to 20X+ slower than accessing a register
- Strategy: Keep variables in registers as much as possible

# Using Aarch32 Registers

- There are two basic groups of registers, **general purpose** and **special use**
- **General purpose registers** can be used to contain up to 32-bits of data, but you must follow the **rules** for their use
  - Rules specify how registers are to be used so software can communicate and share the use of registers (later slides)
- **Special purpose registers:** dedicated hardware use (like r15 the pc) or **special use** when used with certain instructions (like r13 & r14)
- r15/pc is the program counter that contains the address of an instruction being executed (not exactly ... later)

Special Use Registers  
program counter

r15/pc

Special Use Registers  
function call implementation  
& long branching

r14/lr  
r13/sp  
r12/ip  
r11/fp

Preserved registers  
Called functions **can't change**

r10  
r9  
r8  
r7  
r6  
r5  
r4

Scratch Registers  
First 4 Function Parameters  
Function return value  
Called functions **can change**

r3  
r2  
r1  
r0

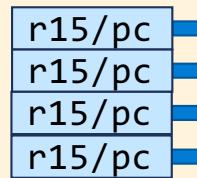
# How r1 (pc) is used: CPU Operational Overview

**Everything** has a **memory address**: **instructions & data**

- Machine code uses addresses for loops, branches, function calls, variables, etc.

## 1 Fetch

- read the instruction into memory (**fetch**)
  - program counter is automatically incremented (+4) to contain the address of the next instruction in memory
- Instructions are 32 bits



## 2 Decode

- Decodes the instruction** and sets up execution

## 3 Execute

- CPU completes the **execution** of the instruction
- Execution may alter the pc to take branches, etc.
- Go to fetch**

text segment in memory		corresponding assembly
address	contents	
0001042c <inloop>:	1042c: e3530061	cmp r3, 0x61
	10430: ba000002	blt 10440 <store>
	10434: e353007a	cmp r3, 0x7a
	10438: ca000000	bgt 10440 <store>
	1043c: e2433020	sub r3, r3, #32
00010440 <store>:	10440: e7c13002	strb r3, [r1, r2]
	10444: e2822001	add r2, r2, 0x1
	10448: e7d03002	ldrb r3, [r0, r2]
	1044c: e3530000	cmp r3, 0x0
	10450: 1affffff	bne 1042c <inloop>

# AArch32 Instruction Categories

- **Data movement to/from memory**
  - Data Transfer Instructions between memory and registers
    - Load, Store
- **Arithmetic and logic**
  - Data processing Instructions (registers only)
    - Add, Subtract, Multiply, Shift, Rotate, ...
- **Control Flow**
  - Compare, Test, If-then, Branch, function calls
- **Miscellaneous**
  - Traps (OS system calls), Breakpoints, wait for events, interrupt enable/disable, data memory barrier, data synchronization barrier
  - Many others that we will not cover in the class

Arithmetic and logic

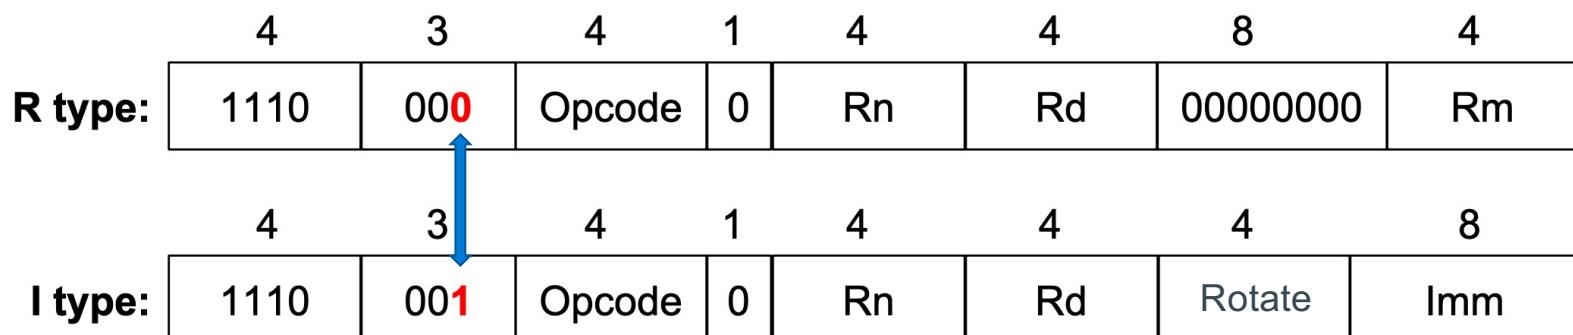
Data Movement

Control Flow

Miscellaneous

## Basic Arm Instructions

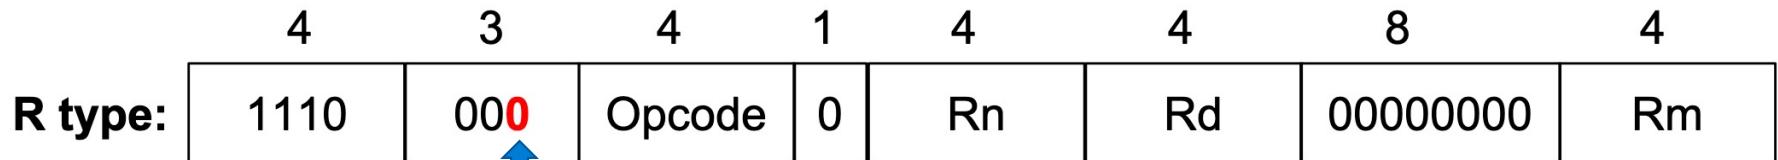
- Instructions consist of several fields that **encode** the **opcode** and arguments to the **opcode**
- Special fields enable extended functionality - later
- Several 4-bit **operand** fields for specifying the **source and destination** of the operation, usually one of the 16 registers
- Embedded constants (*"immediate values"*) of various size and "configuration"
- Basic Data processing instruction formats (below)



## R (register) -Type Data Processing

- Instructions that process data using three-register arguments
- The general instruction format is (not all fields will be in every instruction)

opcode      Rd (destination), Rn (operand 1), Rm (operand 2)



add r0, r1, r3

is encoded as

1110 0000 1000 0001 0000 0000 0000 0011

in hex is

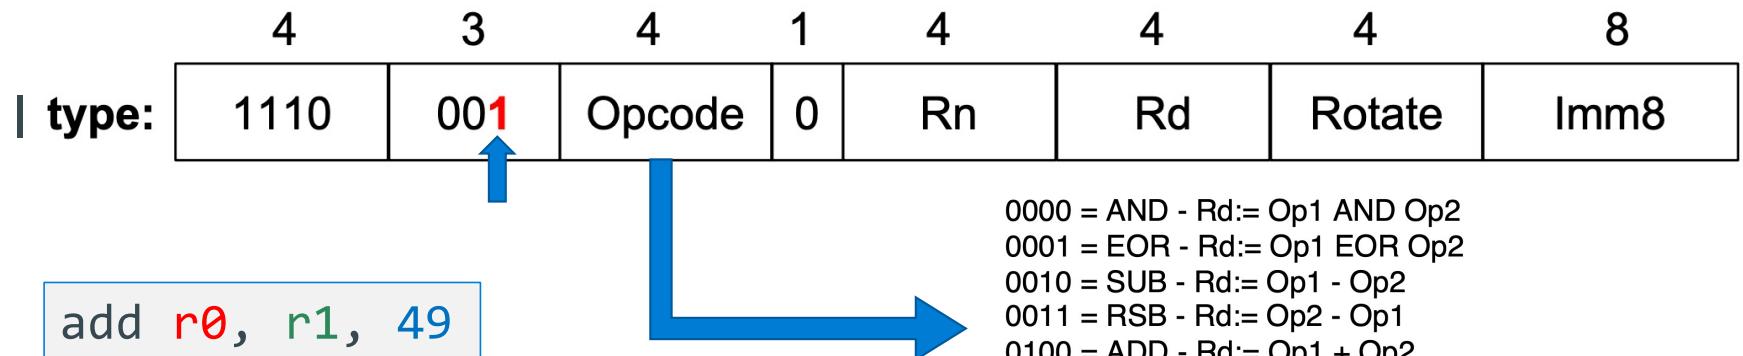
0xe0810003

0000 = AND - Rd:= Op1 AND Op2  
0001 = EOR - Rd:= Op1 EOR Op2  
0010 = SUB - Rd:= Op1 - Op2  
0011 = RSB - Rd:= Op2 - Op1  
0100 = ADD - Rd:= Op1 + Op2  
0101 = ADC - Rd:= Op1 + Op2 + C  
0110 = SBC - Rd:= Op1 - Op2 + C - 1  
0111 = RSC - Rd:= Op2 - Op1 + C - 1  
1000 = TST - set condition codes on Op1 AND Op2  
1001 = TEQ - set condition codes on Op1 EOR Op2  
1010 = CMP - set condition codes on Op1 - Op2  
1011 = CMN - set condition codes on Op1 + Op2  
1100 = ORR - Rd:= Op1 OR Op2  
1101 = MOV - Rd:= Op2  
1110 = BIC - Rd:= Op1 AND NOT Op2  
1111 = MVN - Rd:= NOT Op2

# I (immediate) - Type Data Processing

- Instructions that process data using two registers and a constant (in the instruction)
- The general instruction format is (not all fields will be in every instruction)

opcode    Rd (destination), Rn (operand 1), constant



add r0, r1, 49

is encoded as

1110 0000**0** 1000 **0001** **0000** 0000 **0011** **0001**

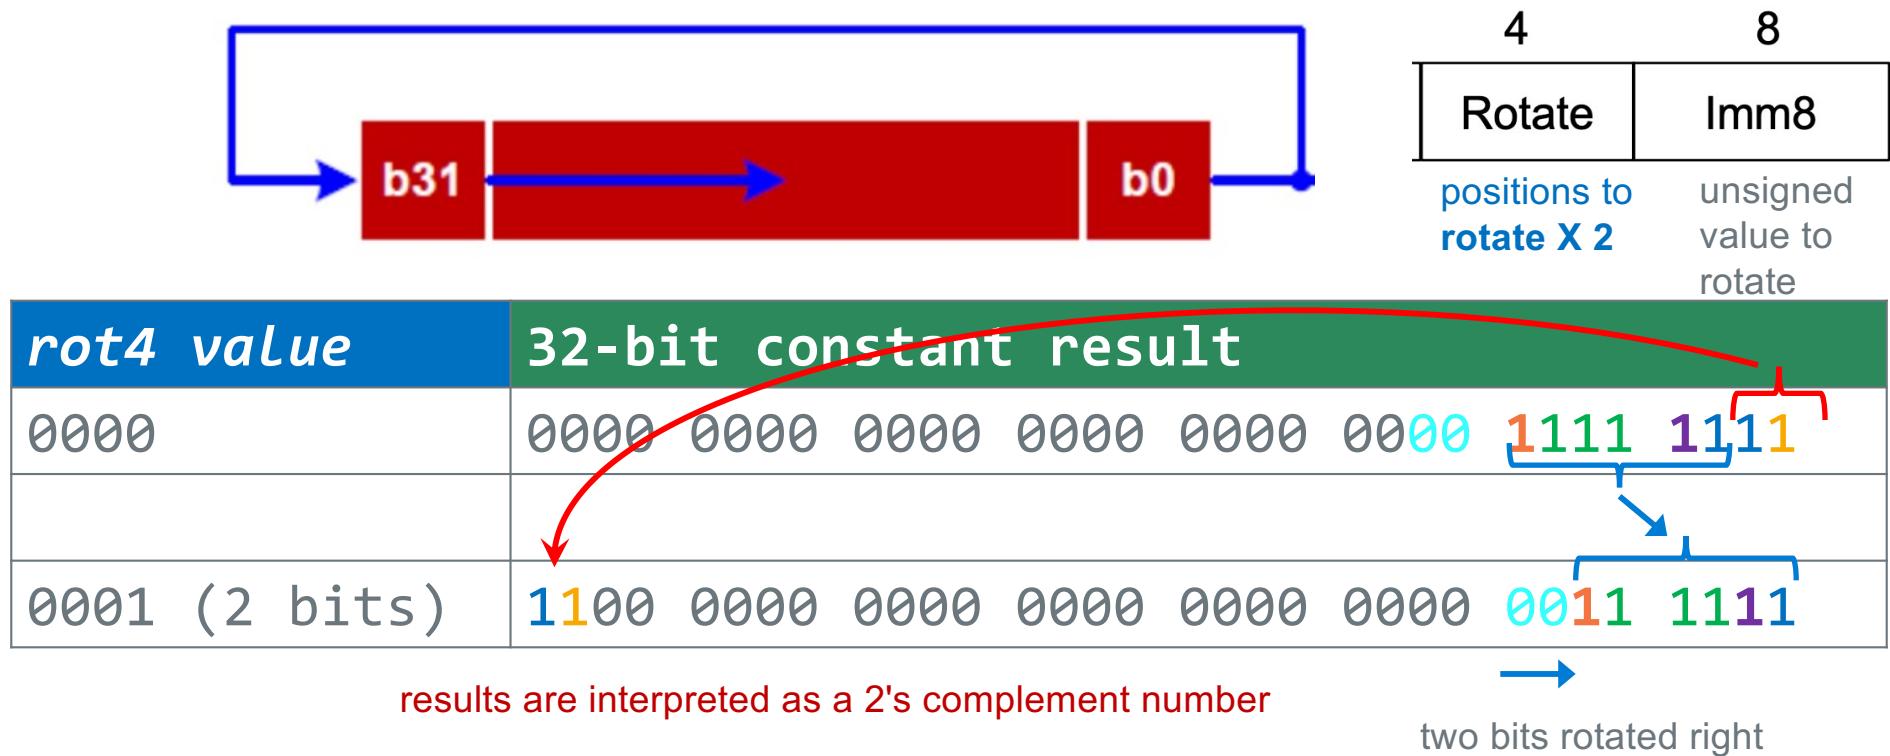
in hex is

0xe08**1**003**1**

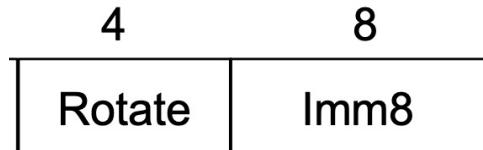
0000 = AND - Rd:= Op1 AND Op2
0001 = EOR - Rd:= Op1 EOR Op2
0010 = SUB - Rd:= Op1 - Op2
0011 = RSB - Rd:= Op2 - Op1
0100 = ADD - Rd:= Op1 + Op2
0101 = ADC - Rd:= Op1 + Op2 + C
0110 = SBC - Rd:= Op1 - Op2 + C - 1
0111 = RSC - Rd:= Op2 - Op1 + C - 1
1000 = TST - set condition codes on Op1 AND Op2
1001 = TEQ - set condition codes on Op1 EOR Op2
1010 = CMP - set condition codes on Op1 - Op2
1011 = CMN - set condition codes on Op1 + Op2
1100 = ORR - Rd:= Op1 OR Op2
1101 = MOV - Rd:= Op2
1110 = BIC - Rd:= Op1 AND NOT Op2
1111 = MVN - Rd:= NOT Op2

## How are I – Type Constants Encoded in the instruction?

- AArch32 provides only 8-bits for specifying an immediate constant value
- This does not provide a very large range of constant values (0-255)
- Imm8 expands to 32 bits and does a rotate right to achieve additional constant values



## Rot4 - Imm8 Values



positions to rotate X 2      unsigned value to rotate

- How would **256** be encoded?
  - rotate = 12, imm8 = 1
- **Bottom line:** the assembler will do this for you
- If you try and use an immediate value that it cannot generate it will give an error
- There is a workaround - later

Rotate	Bits Used	Range
0		0 - 255
1		-2147483648 - 1073741887
2		-2147483648 - 1879048207
3		-2147483648 - 2080374787
4		-2147483648 - 2130706432
5		4194304 - 1069547520
6		1048576 - 267386880
7		262144 - 66846720
8		65536 - 16711680
9		16384 - 4177920
10		4096 - 1044480
11		1024 - 261120
12	00000001	256 - 65280
13		64 - 16320
14		16 - 4080
15		4 - 1020

results are interpreted as a 2's complement number

Version 1.07

# UCSD CSE 30

## Computer Organization and Systems Programming

### Aarch32 Assembly - Introduction

Week 6

Lecture 18

Keith Muller



## First Look: Copying Values To Registers - MOV

```
mov r0, r1
```

// Copies all 32 bits  
// of the value held  
// in register r1 into  
// the register r0

register r1



register r0

```
mov r0, 100
```

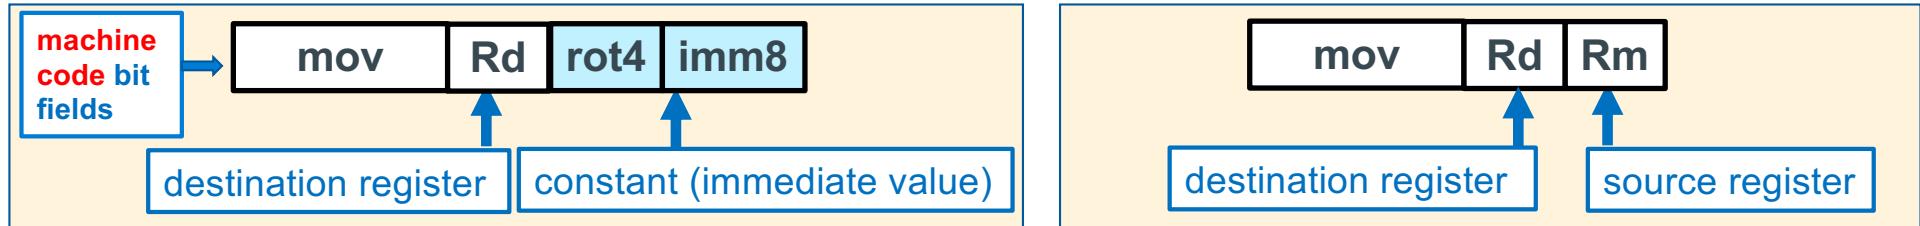
// Expands an imm8 value 100  
// stored in the instruction  
// into the register r0

100



register r0

## mov – Copies Register Content between registers



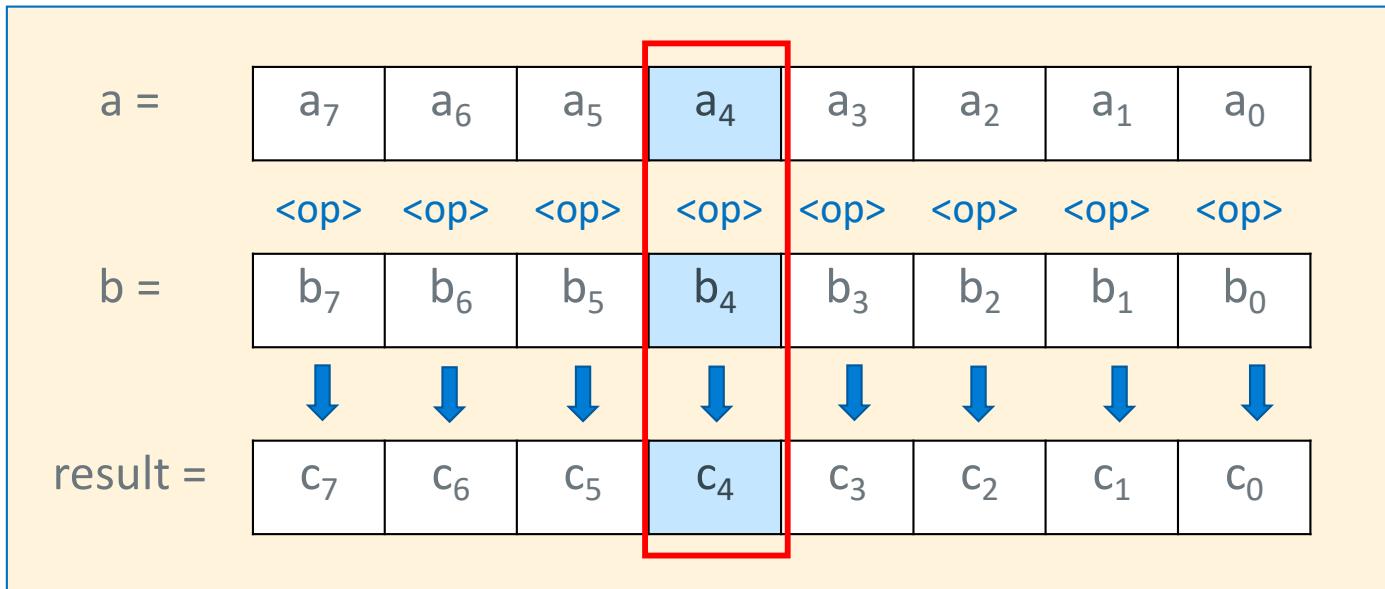
R type:	4	3	4	1	4	4	8	4
	1110	000	Opcode	S	0000	Rd	00000000	Rm
I type:	1110	001	Opcode	S	0000	Rd	Rotate	Imm8

↳ {  
  1101 - MOV  
  1111 - MVN}

```
mov Rd, constant // Rd = constant
mov Rd, Rm // Rd = Rm
```

```
mov r1, r5 // r1 = r5
mov r1, 1 // r1 = 1
mov r1, -4 // r1 = -4
```

# What is a Bitwise Operation?



- Bitwise operators are applied to each of the corresponding bit positions in each variable
- Each bit position of the result depends only on bits in the **same bit position** within the operands

## Bitwise Not (vs Boolean Not)

```
in C  
int output = ~a;
```

a	$\sim a$
0	1
1	0

Bitwise NOT  
 $\sim$  1100  
---  
0011

	Bitwise Not							
number	0101	1010	0101	1010	1111	0000	1001	0110
$\sim$ number	1010	0101	1010	0101	0000	1111	0110	1001

<i>Meaning</i>	<i>Operator</i>	<i>Operator</i>	<i>Meaning</i>
Boolean NOT	$!b$	$\sim b$	Bitwise NOT

Boolean operators act on the entire value not the individual bits

Type	Operation	result
bitwise	$\sim 0x01$	1111 1111 1111 1111 1111 1111 1111 1111 1110
Boolean	$!0x01$	0000 0000 0000 0000 0000 0000 0000 0000 0000

## First Look: Copying Values To Registers – MVN (negate)

`mvn r0, r1`

// Copies all 32 bits  
// of the value held  
// in register r1 into  
// the register r0  
// then does a bitwise NOT

register r1



register r0

`mvn r0, 12`

// Expands an imm8 value 0x0c  
// stored in the instruction  
// into a register then does  
// a bitwise NOT

register r0

0x0c



0xffff fff3

Bitwise NOT

~ 1100  
---  
0011

- A **bitwise NOT** operation

0x 0c

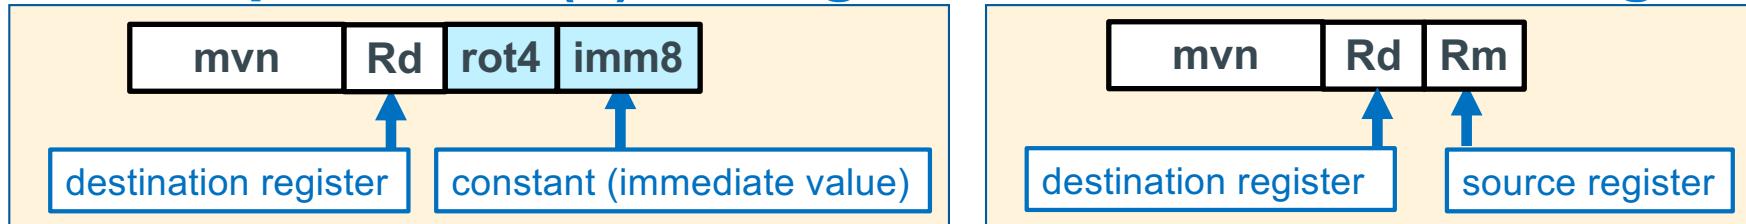
↓ imm8 expansion

0x0000000c

↓ bitwise not

0xffffffff3

## mvn – Copies NOT (~) of Register content between registers

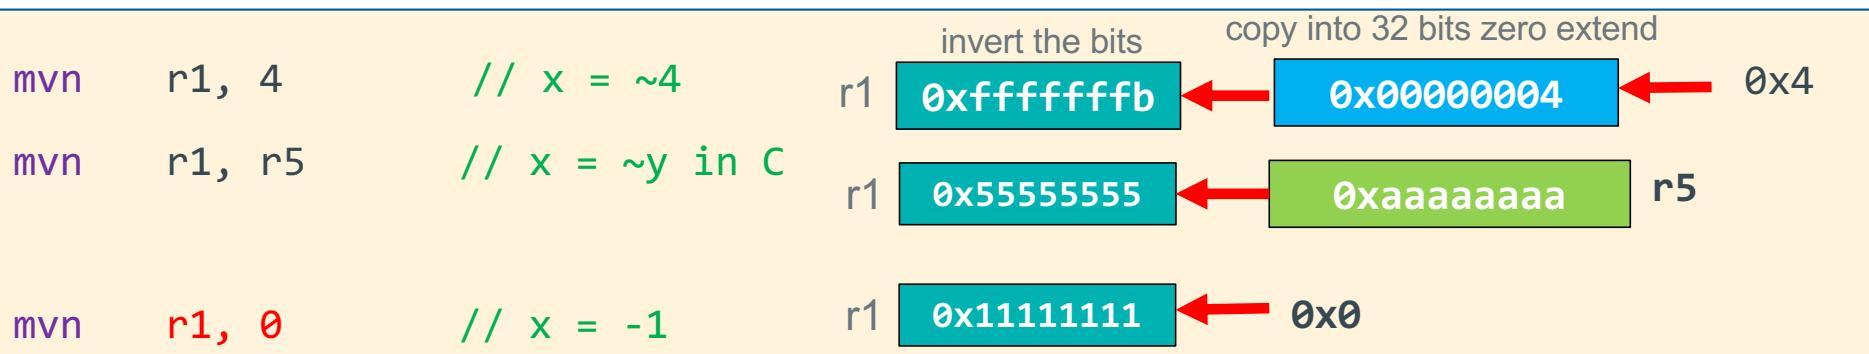


```
mvn Rd, constant // Rd = constant
mvn Rd, Rm // Rd = Rm
```

Bitwise NOT  
 $\sim \begin{matrix} 1 & 1 & 0 & 0 \\ \text{---} \\ 0 & 0 & 1 & 1 \end{matrix}$

**bitwise NOT** operation. Immediate (constant) version copies to 32-bit register, then does a bitwise NOT

imm8	extended imm8	inverted imm8	signed base 10
0x00	0x00 00 00 00	0xff ff ff ff	-1
0xff	0x00 00 00 ff	0xff ff ff 00	-256



## First Look: Add/Sub Registers

add r0, r1, r2 register r1 + register r2

// Adds r1 to r2 and  
// stores the result  
// in r0

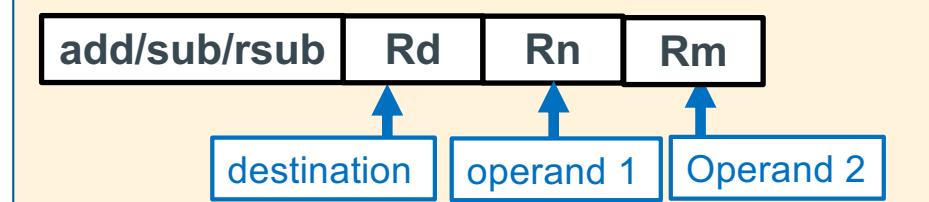
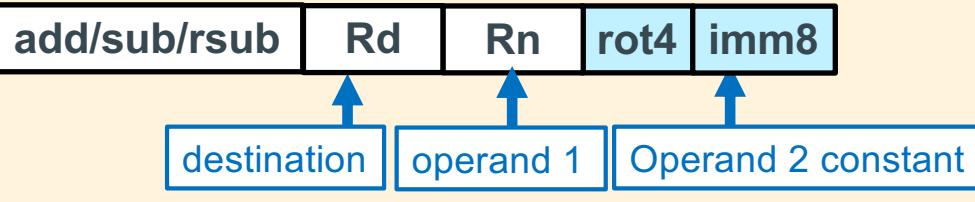
register r0

sub r0, r1, 100 register r1 - 100

// Perform r1 - 100 and  
// stores the result in  
// r0

register r0

## add/sub/rsub – Add or Subtract two integers



```

add Rd, Rn, constant      // Rd = Rn + constant
sub Rd, Rn, constant      // Rd = Rn - constant
rsub Rd, Rn, constant     // Rd = constant - Rn
add Rd, Rn, Rm            // Rd = Rn + Rm
sub Rd, Rn, Rm            // Rd = Rn - Rm
rsub Rd, Rn, Rm           // Rd = Rm - Rn
  
```

```

mov r5, 5      // r5 = 5
mov r7, 7      // r7 = 7
add r7, r7, r5 // r7 = 12 r5 = 5
  
```

```

add r1, r2, r3 // r1 = r2 + r3
sub r1, r1, 1   // r1 = r1 - 1; or r1--
add r1, r2, 234 // r1 = r2 + 234
  
```

## Writing a Sequence of Add & Subtract Instructions

- You need to perform the following sequence of integer adds/subtracts
$$a = b + c + d - e;$$
- Since ARM uses a three-operand instruction set, you can only operate on two operands at a time
- So, you need to use one register as an **accumulator** and create a **sequence of add instructions** to build up the solution

```
r0 ← a  
r1 ← b  
r2 ← c  
r3 ← d  
r4 ← e
```

```
a = b + c + d - e;  
r0 = r1 + r2 + r3 - r4;  
r0 = ((r1 + r2) + r3) - r4;  
r0 = r1 + r2;  
r0 = r0 + r3  
r0 = r0 - r4
```

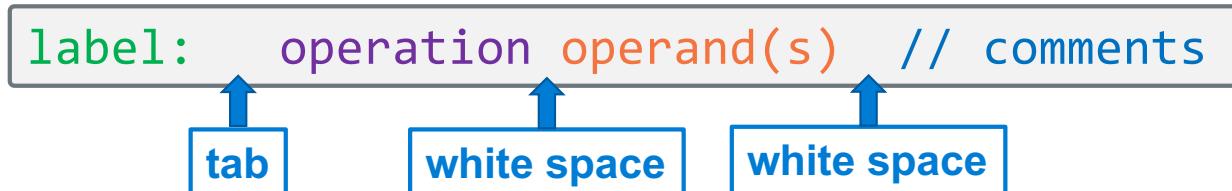
```
add r0, r1, r2  
add r0, r0, r3  
sub r0, r0, r4
```

```
a = (b + c) - 5;  
r0 = (r1 + r2) - 5;
```

```
add r0, r1, r2  
sub r0, r0, 5
```

## Overview: Line Layout in an Arm Assembly Source File - 1

column 1      column 2      column 3      column 4



- Assembly language source text files are **line oriented** (each ending in a '\n')
- **Each line represents** a **starting address in memory** and does **one of**:
  1. Specifies the **contents of memory** for a **variable** (segments containing data)
  2. Specifies the **contents of memory** for an **instruction** (text segment)
  3. **Assembler directives** tell the assembler to do something (for example, change label scope, define a macro, etc.) that **does not allocate memory**
- **Each line is organized into up to four columns**
  - Not every column is used on each line
  - Not every line will result in memory being allocated

## Overview: Line Layout in an Arm Assembly Source File - 2

```
label: operation operand(s) // comment  
          // assembler directive below  
cnt:   .word 5           /* define a global int cnt = 5;  
          /* instruction below */  
        add    r1    r2, r3      // add the values
```

1. **Labels** (optional); starts in column 1
  - **Only used** when you need to **associate a name to a starting location in memory**; You can then refer to the address **by name** in an **instruction**
2. **Operation type 1: assembler directives** (all start with a period e.g. `.word` )
3. **Operation Type 2: assembly language instructions**
4. **Zero or more operands** as required by the instruction or assembler directive
5. **Comments C and C++ style**; also @ in the place of a C++ comment //

# Labels in Arm Assembly

local label

```
.Lmesg: .asciz "Hello CSE30! We Are CountinG UpPpER cASE letters!"
```

... label .Lmesg is the starting address for the ascii string

Regular label

```
main:
```

label main is the starting address of the push instruction

```
push {fp, lr}
```

local label

```
...
```

.Lwhile: label .Lwhile is the starting address of the add instruction

```
add r2, r2, 1 // increment char pointer
```

- Remember, a Label associates a name with memory location
- Regular Label:**
  - Used with a Function name (label) or for static variables in any of the data segments
- Local Label:** Name starts with `.L` (local label prefix) only usable in the same file
  - Targets for branches (if), switch, goto, break, continue, loops (for, while, do-while)
  - Anonymous variables (string not foo in `char *foo = "anonymous variable"`)
  - Read only literals when allocated in the text segment – special case)

## Assembler Directives: Label Scope Control (Normal Labels only)

```
.extern printf  
.extern fgets  
.extern strcpy  
.global fbuf
```

`.extern <label>`

- **Imports** `label` (function name, symbol or a static variable name);
- An address associated with the label from another file can be used by code in this file

`.global <label>`

- **Exports** `label` (or `symbol`) to be visible outside the source file boundary (other assembly or c source)
- `label` is either a `function name` or a `global variable name`
- **Without** `.global`, **by default** labels are local to the file from the point where they are defined

## Assembler Directives: .equ and .equiv

```
.equ    BLKSZ, 10240      // buffer size in bytes
.equ    BUFCNT, 100*4      // buffer for 100 ints
.equiv  STRSZ, 128        // buffer for 128 bytes
.equiv  STRSZ, 1280       // ERROR! already defined!
.equ    BLKSZ, STRSZ * 4   // redefine BLKSZ from here
```

**.equ <symbol>, <expression>**

- Defines and sets the value of a symbol to the evaluation of the expression
- Used for specifying constants, like a `#define` in C
- You can (re)set a symbol many times in the file, last one seen applies

```
.equ    BLKSZ, 10240      // buffer size in bytes
// other lines
.equ    BLKSZ, 1024        // buffer size in bytes
```

**.equiv <symbol>, <expression>**

- .equiv directive is like .equ except that the assembler will signal an error if symbol is already defined

## Example: Assembler Directive and Instructions

