

Version 1.09

UCSD CSE 30


Aarch32 Assembly – Stack Frames, Bit Ops, Shifts

Week 9

Lecture 25

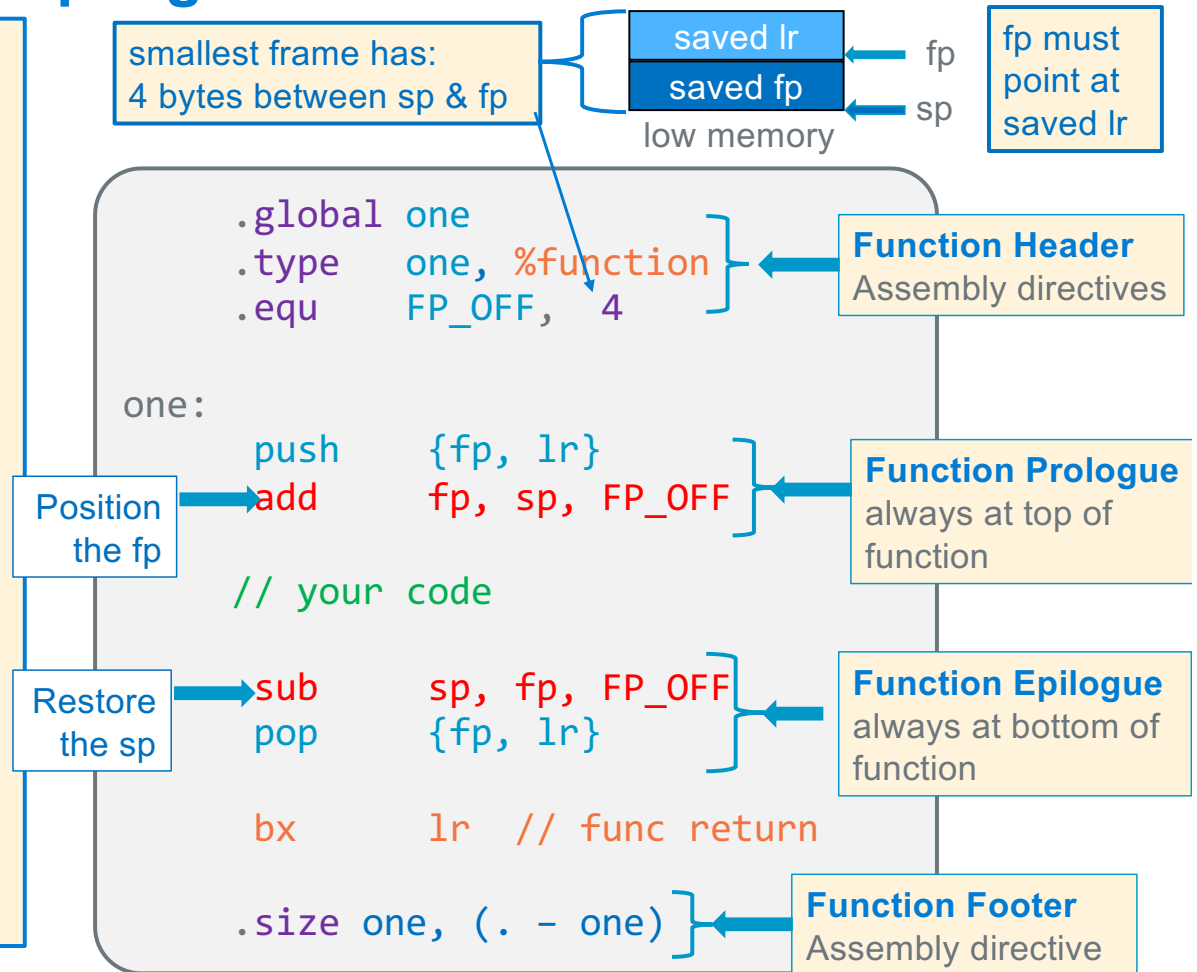
Keith Muller

CDC 3600 -1963



Function Prologue and Epilogue: Minimum Stack Frame

- **Function prologue** creates stack frame
 1. push/save registers (**lr** & **fp** minimum) on stack
 2. set **fp** (**add fp, ...**) to point at the saved lr as required for use by this function (later)
- **Function epilogue** removes stack frame
 1. set **sp** to where it was at the push (we may have **moved sp** to allocate space, later slides)
 2. pop/restore registers (**lr** & **fp** minimum) from stack
- In this example **fp** is 4 bytes from **sp**, (**FP_OFF**) but this will vary...




Stack Creation Overview

1. Calculate how much additional space is needed by local variables
2. **After the push, Subtract from the sp the required byte count (+ padding - later slides)**
3. If the variable has an initial value specified: **add code to set the initial value**
 - a) mov and str are useful for initializing simple variables
 - b) loops of mov and str for arrays

```
#define BFSZ 256
int main(void)
{
    char buf[BFSZ]; // BFSZ bytes
    ...
}
```

stack after allocating local space After
sub sp, sp, BFSZ



```
.equ    FP_OFF, 4
.equ    BFSZ, 256
```


main:

**Function
Prologue
Extended**

```
push    {fp, lr}
add     fp, sp, FP_OFF
sub     sp, sp, BFSZ
```

allocate
space for
buf[256]

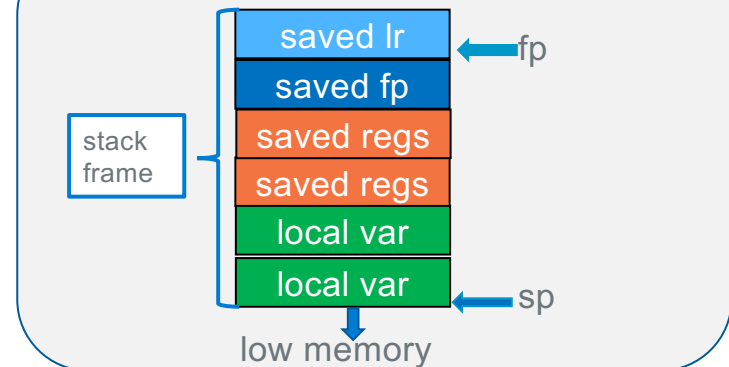
Why is there a `sub, fp, FP_OFF` ?



```
push    {fp, lr}
add     fp, sp, FP_OFF
```

- As you will see, we will move the `sp` to allocate space on the stack for local variables and parameters, so for the `pop` to restore the registers correctly:
- `sp` must point at the last saved preserved register put on the stack by the save register operation: the `push`

So we can add space for local variables!



```
.equ     FRMSZ, 8
push     {fp, lr}
add      fp, sp, FP_OFF
sub      sp, sp, FRMSZ
// your code

sub      sp, fp, FP_OFF
pop      {fp, lr}

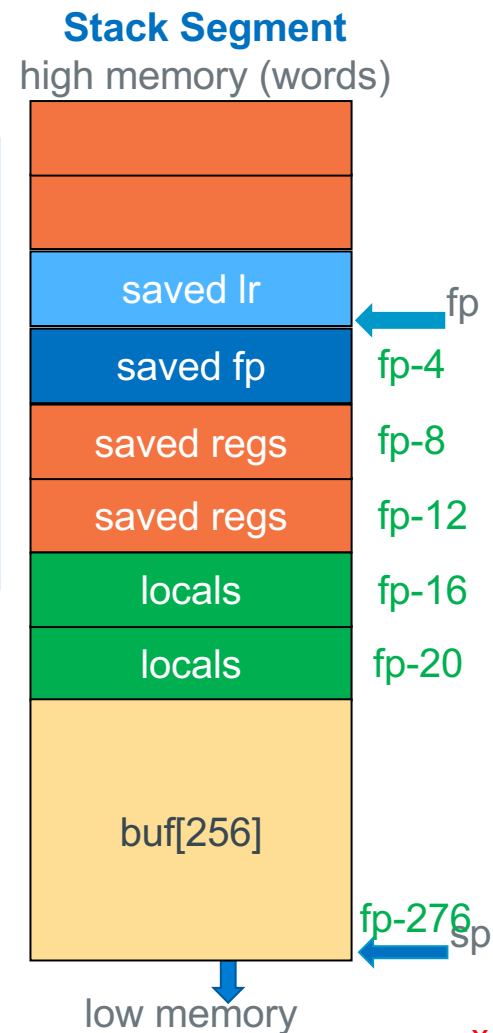
bx       lr // func return
```

- force the `sp` (using the `fp`) to contain the same address it had after the push operation
`sub sp, fp, FP_OFF`

Accessing the Stack Variables Overview

- Access data stored in the stack use `ldr/str` instructions
- Use base register `fp` with offset addressing (either register offset or immediate offset)
- No matter where the stack frame starts on the stack, `fp` always points at the same place in every stack (points at saved `lr`)
- *"hand calculated offset constants and sizes" like -16 and -20 to access items is easy to get wrong, there is an easier way!*

```
.equ  BFSZ, 256           // char array
ldr   r0, [fp, -16]
str   r3, [fp, -20]
sub   r0, fp, 256+20      // r0 = &(buf[0]);
ldr   r1, [r0]            // r1 = buf[0];
str   r3, [r0, 2]         // buf[2] = r3;
```



Variable Alignment on Stack

integer/pointer		short	char
4 bytes		2 bytes	1
Variable Type/Size	Address ends in		
8-bit char -1 byte	0b..0 or 0b..1		
16-bit int -2 bytes	0b.. 0		
32-bit int -4 bytes	0b.. 00		
32-bit pointer -4 bytes	0b.. 00		

- Starting **address alignment requirements** for local variables **stored on the stack** is just like static variables
- sp** must be aligned to **8-bytes** at function entry & exit
 - contents of sp always ends in 0b..**000** at function entry
- Approach we will take (also what compilers often do): allocate all the local variable space as part of the function prologue**
 - Aside: You cannot use .align as assembly directives are for fixed address

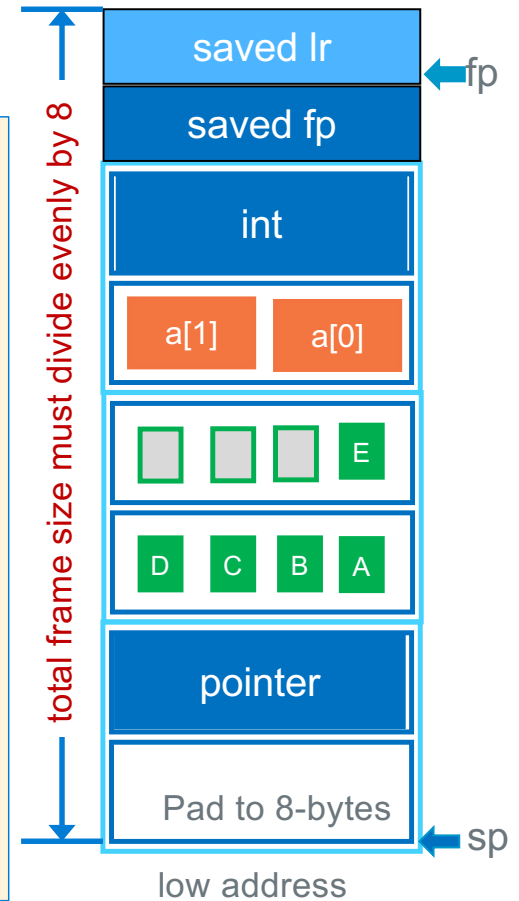
Starting address by size

4 bytes	2 Bytes	1 Byte	Addr. (hex)
Addr = 0x0C	Addr = 0x0E		0x..0F
	Addr = 0x0C		0x..0E
Addr = 0x08	Addr = 0x0A		0x..0D
	Addr = 0x08		0x..0C
Addr = 0x04	Addr = 0x06		0x..0B
	Addr = 0x04		0x..0A
Addr = 0x00	Addr = 0x02		0x..09
	Addr = 0x00		0x..08
			0x..07
			0x..06
			0x..05
			0x..04
			0x..03
			0x..02
			0x..01
			0x..00

Overview: Stack Frame Alignment Rules



- Goal: minimize stack frame size
- Arrays start at a 4-byte boundary (even arrays with only 1 element)
 - Exception: double arrays [] start at an 8-byte boundary
 - struct arrays are aligned to the requirements of largest member
- Space padding when necessary is added at the high address end of a variables allocated space, so the next variable is aligned
- Single chars (and shorts) can be grouped together in same 4-byte word (following the alignment for the short)
- After all the variables have been allocated, add padding at stack frame bottom (low memory) so the total stack frame size (including all saved registers) is a multiple of 8 when the prologue is finished



Stack Frame Design – Step 1 Listing the Local Variables

```
int func(void)
{
    char str[] = "Hi";
    char *ptr = str;
    short buf[3];
    // other code
    int n = 0;
    // other code

    return EXIT_SUCCESS;
}
```

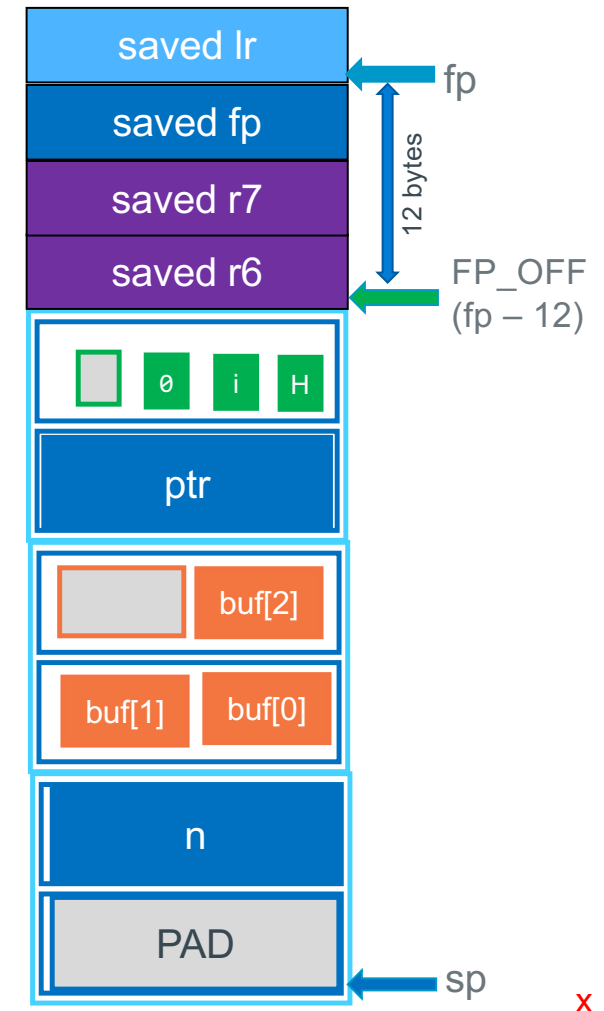
Variable name	Initial Value	Size bytes	Alignment pad to next	Total Size
char str[]	"Hi"	3	1	4
char *ptr	str	4	0	4
short buf[3]		3 * 2	2	8
int n	0	4	0	4
<sub total>				20

- Create a table of **all the variables defined throughout the entire function** starting from function start to the end of the function
- **For each variable:** list its size, initial value if any, alignment padding and sum to total size
 - **When needed:** padding after the variable (the high address side) to fill out the allocation

Stack Frame Design – Step 2 Layout the Frame & Size It

		4 bytes	2 bytes	1	
Variable name	Initial Value	Size bytes	Alignment pad to next	Total Size	
char str[]	"Hi"	3	1	4	
char *ptr	str	4	0	4	
short buf[3]		3 * 2	2	8	
int n	0	4	0	4	
<sub total>				20	


Allocation Type	Total
$FP_OFF + 4 = 12 + 4 = 16$	16
Local Variables Sub total	20
Space for parameters on stack (later)	0
total before pad	36
Pad as needed to align 8-byte boundary	4
TOTAL Size for entire frame	40



x

Stack Frame Design – Step 3 Generate the Offsets from fp

- **Word offset** is a way to visualize the distance from fp for calculating offset values
- Better to have the assembler to generate readable offsets for use with **str** and **ldr**
 1. Easy to add and remove variable allocations from the design
 2. Creates well documented names for each variable: **ldr r0, [fp -20]** is hard to read
 3. Automatically calculates the total size of the stack frame used by local variables



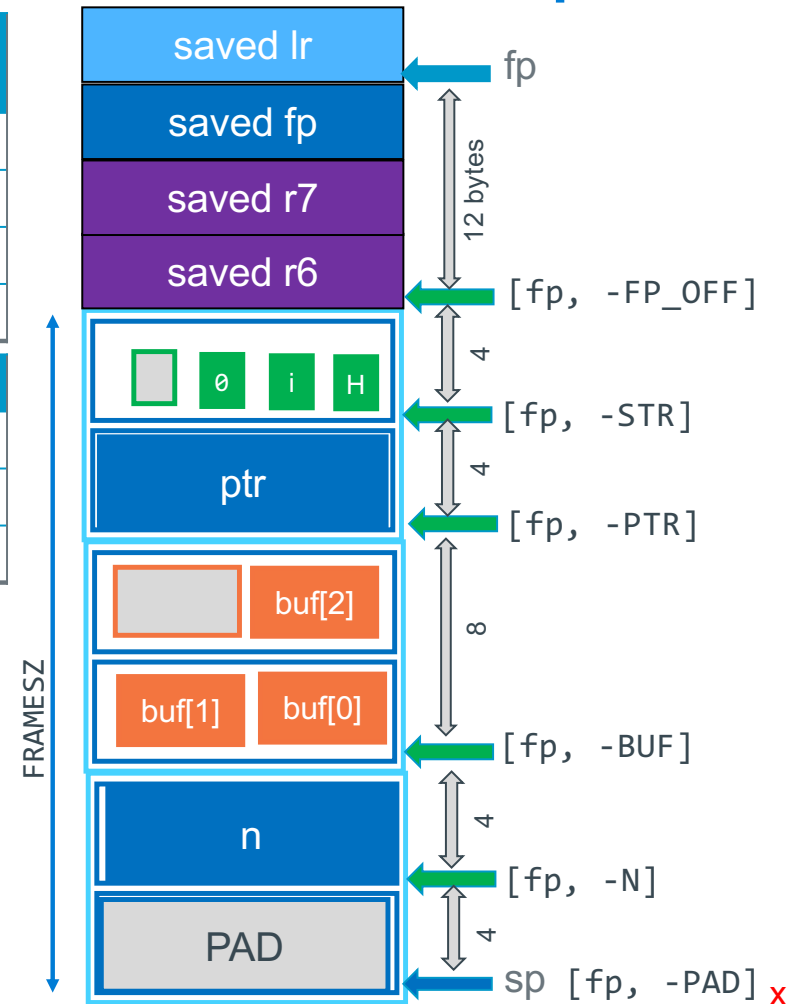
Stack Frame Design – Step 3 Generate the offsets from fp

Variable name	Initial Value	Size bytes	Alignment pad to next	Total Size
char str[]	"Hi"	3	1	4
char *ptr	str	4	0	4
short buf[3]		3 * 2	2	8
int n	0	4	0	4

Allocation Type	Total
FP_OFF + 4 = 12 + 4 = 16	16
Pad to get to 8-byte boundary	4
FRAMESZ space for Locals (PAD - FP_OFF)	24

Distance
Offsets
from fp

```
.equ FP_OFF, 12 // local base
// NAME, SIZE + prev_name
.equ STR, 4 + FP_OFF
.equ PTR, 4 + STR
.equ BUF, 8 + PTR
.equ N, 4 + BUF
.equ PAD, 4 + N
.equ FRAMESZ PAD - FP_OFF
```



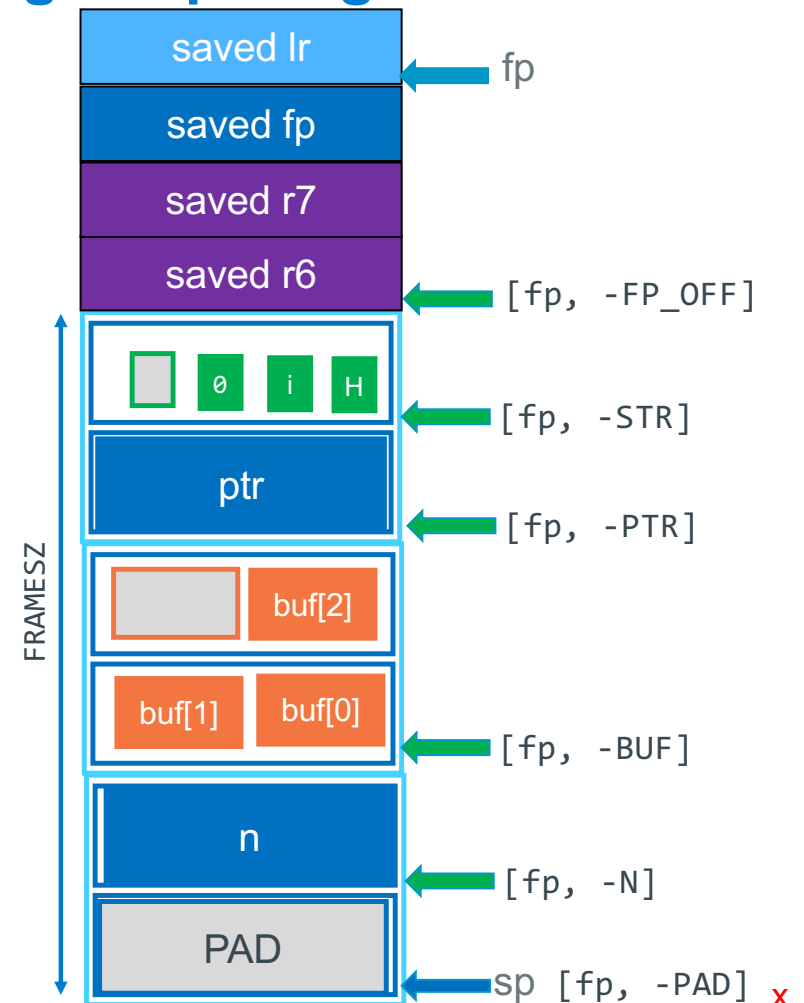
Stack Frame Design – Step 4 Modifying the prologue

Distance
Offsets
from fp

```
.equ  FP_OFF,    12  // local base
      // NAME,    SIZE + prev_name
.equ  STR,       4 + FP_OFF
.equ  PTR,       4 + STR
.equ  BUF,       8 + PTR
.equ  N,         4 + BUF
.equ  PAD,       4 + N
.equ  FRAMESZ    PAD - FP_OFF
```

```
main:
  push    {r6, r7, fp, lr}
  add     fp, sp, FP_OFF
  sub     sp, sp, FRAMESZ // add for locals
  → // no change to epilogue ←
```

variable	arm ldr/str statement examples
n	ldr/str r0, [fp, -N]
buf[1]	ldrh/strh r0, [fp, -BUF + 2]
&(str[0])	sub r0, fp, STR



Stack Frame Design – Step 5 Initialize the variables


```
char str[] = "Hi";
char *ptr = str;
short buf[3];
// other code
int n = 0;
// other code
```

main:

```
push    {r6, r7, fp, lr}
add     fp, sp, FP_OFF
sub     sp, sp, FRAMESZ
sub     r6, fp, STR    // &(str[0])
str     r6, [fp, -PTR]
mov     r6, 'H'
strb    r6, [fp, -STR]
mov     r6, 'i'
strb    r6, [fp, -STR+1]
mov     r6, 0
strb    r6, [fp, -STR+2]
// other code
mov     r6, 0
str     r6, [fp, -N]
```

Used
in PA5

FRAMESZ



Version 1.09

UCSD CSE 30


Aarch32 Assembly – Stack Frames, Bit Ops, Shifts

Week 9

Lecture 26

Keith Muller

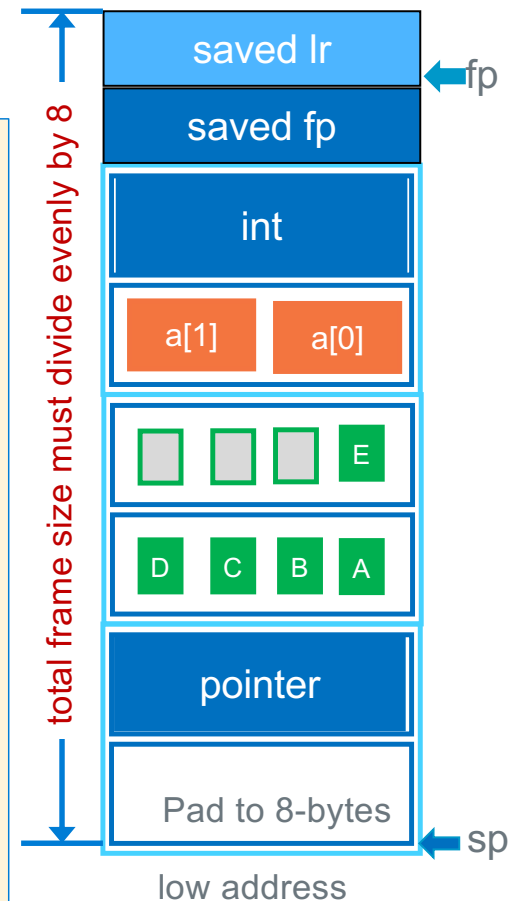
Cray XMP - 1992



Overview: Stack Frame Alignment Rules



- Goal: minimize stack frame size
- Arrays start at a 4-byte boundary (even arrays with only 1 element)
 - Exception: double arrays [] start at an 8-byte boundary
 - struct arrays are aligned to the requirements of largest member
- Space padding when necessary is added at the high address end of a variables allocated space, so the next variable is aligned
- Single chars (and shorts) can be grouped together in same 4-byte word (following the alignment for the short)
- After all the variables have been allocated, add padding at stack frame bottom (low memory) so the total stack frame size (including all saved registers) is a multiple of 8 when the prologue is finished




Local Variables: Stack Frame Design Practice

Example shows allocation **without reordering** variables to optimize space

```
short a[3];
short *ptr1;
char tmp;
char *ptr2;
char nm[] = "cse30";
```

```
.equ  FP_OFF,      4  // Local base
// NAME,          SIZE + prev_name
.equ  A,           8 + FP_OFF
.equ  PTR1,        4 + A
.equ  TMP,         4 + PTR1
.equ  PTR2,        4 + TMP
.equ  NM,          8 + PTR2
.equ  PAD,         4 + NM
.equ  FRAMESZ     PAD - FP_OFF // for locals
```



When writing real code, you do not have to put all locals on the stack


- Place locals in registers if they fit, are accessed often, and
- You do not need their address (they are not an output variable in a function call)

Local Variables: Stack Frame Design Reordering

Example shows allocation **with reordering** variables to optimize space

```
short a[3];
short *ptr1;
char *ptr2;
char tmp;
char nm[] = "cse30";
```


```
.equ  FP_OFF,      4  // Local base
// NAME,          SIZE + prev_name
.equ  A,           8 + FP_OFF
.equ  PTR1,        4 + A
.equ  PTR2,        4 + PTR1
.equ  TMP,         size 2 + PTR2
.equ  NM,          size change 6 + TMP
.equ  PAD,         0 + NM // not needed
.equ  FRAMESZ     PAD - FP_OFF
```



When writing real code, you do not have to put all locals on the stack

- Place locals in registers if they fit, are accessed often, and
- You do not need their address (they are not an output variable in a function call)

Entire source file



	<i>Evaluated into r0</i>
<code>&(a[1])</code>	<code>sub r0, fp, A - 2</code>
<code>&(a[1])</code>	<code>add r0, fp, -A + 2</code>
<code>&(nm[1])</code>	<code>add r0, fp, -NM + 1</code>
<code>ptr2</code>	<code>add r0, fp, -PTR2</code>

```

.arch armv6
.arm
.fpu vfp
.syntax unified
// globals etc here
.text
.type    doit, %function
.global  doit
.equ     EXIT_SUCCESS, 0
.equ     FP_OFF, 4 // Local base
.equ     A, 8 + FP_OFF
.equ     PTR1, 4 + A
.equ     PTR2, 4 + PTR1
.equ     TMP, 2 + PTR2
.equ     NM, 6 + TMP
.equ     PAD, 0 + NM
.equ     FRAMESZ PAD - FP_OFF

doit:
push     {fp, lr}
add      fp, sp, FP_OFF
sub      sp, sp, FRAMESZ
// doit() code goes here
mov      r0, EXIT_SUCCESS

sub      sp, fp, FP_OFF
pop      {fp, lr}
bx       lr

.size    doit, (. - doit)
.section .note.GNU-stack,"",%progbits

.end

```

With large frames you may need to use `ldr` if the immediate value `FRAMESZ` does not fit in `imm8` (`r3` is not a parameter in this example)

```


ldr      r3, =FRAMESZ
sub      sp, sp, r3

```

Passing an Output Parameter

```
#define BUFSZ 256
int main(void)
{
    char buf[BUFSZ];
    if (fgets(buf, BUFSZ, stdin) != NULL)
        printf("%s", buf);
    return EXIT_SUCCESS;
}
```

`char *fgets(char *s, int size, FILE *stream);`
 returns *s or NULL `r0`, `r1`, `r2`



if the immediate value of BUF does not fit in imm8
`ldr r0, =BUF`
`sub r0, fp, r0`

if the immediate value of BUFSZ does not fit in imm8
`ldr r1, =BUFSZ`


```
.extern printf
.extern fgets
.extern stdin ← stdin is a global variable pointer *FILE
.section .rodata
.Lpfstr .string "%s"
.text
// function header stuff not shown
.equ BUFSZ, 256
.equ FP_OFF, 4
.equ BUF, BUFSZ + FP_OFF
.equ FRAMESZ, BUF - FP_OFF

main:
    push {fp, lr}
    add fp, sp, FP_OFF
    sub sp, sp, FRAMESZ
    sub r0, fp, BUF
    mov r1, BUFSZ
    ldr r2, =stdin
    ldr r2, [r2]
    bl fgets
    cmp r0, NULL
    beq .Ldone
    mov r1, r0
    ldr r0, =.Lpfstr
    bl printf
.Ldone: // rest of file not shown
```


`r0 = &(buf[0]);`
`r1 = BUFSZ;`
`r2 = stdin`

Writing Functions: Receiving an Output Parameter - 1

```
#define BFSZ 256
void fillbuf(char *s, int len, char fill);
int main(void)    r0,    r1,    r2
{
    char buf[BFSZ];
    fillbuf(buf, BFSZ, 'A');
    return EXIT_SUCCESS;
}
```




```
void fillbuf(char *s, int len, char fill)
{
    r0,    r1,    r2
    char enptr = s + len;
    while (*s < enptr)
        *(s++) = fill;
}
```



Writing Function: Receiving an Output Parameter - 2

```
void      r0,      r1,      r2
fillbuf(char *s, int len, char fill)
{
    char enptr = s + len;
    while (s < enptr)
        *(s++) = fill;
}
```

Using r1 for endptr



```
fillbuf:
    push    {fp, lr}           // stack frame
    add     fp, sp, FP_OFF     // set fp to base

    add     r1, r1, r0         // copy up to r1 = bufpt + cnt
    cmp     r0, r1             // are there any chars to fill?
    bge     .Ldone             // nope we are done


.Ldowhile:
    strb    r2, [r0]           // store the char in the buffer
    add     r0, 1              // point to next char
    cmp     r0, r1             // have we reached the end?
    blt     .Ldowhile          // if not continue to fill

.Ldone:
    sub     sp, fp, FP_OFF     // restore stack frame top
    pop     {fp, lr}           // restore registers
    bx      lr                 // return to caller
```

Passing More Than Four Arguments - 1

```
r0 = function(r0, r1, r2, r3, arg5, arg6, ... argn)
      arg1, arg2, arg3, arg4, ...
```

- Each argument is a value that must fit in 32-bits
- **Args > 4 are in the caller's stack frame and arg 5 always starts at fp+4**
 - **At the function call (bl) sp points at arg5**
 - Additional args are higher up the stack, with one argument "slot" every 4-bytes
- Called functions have the **right to change stack args** just like they can change the register args!
- Caller must assume all args including ones on the stack are changed by the caller

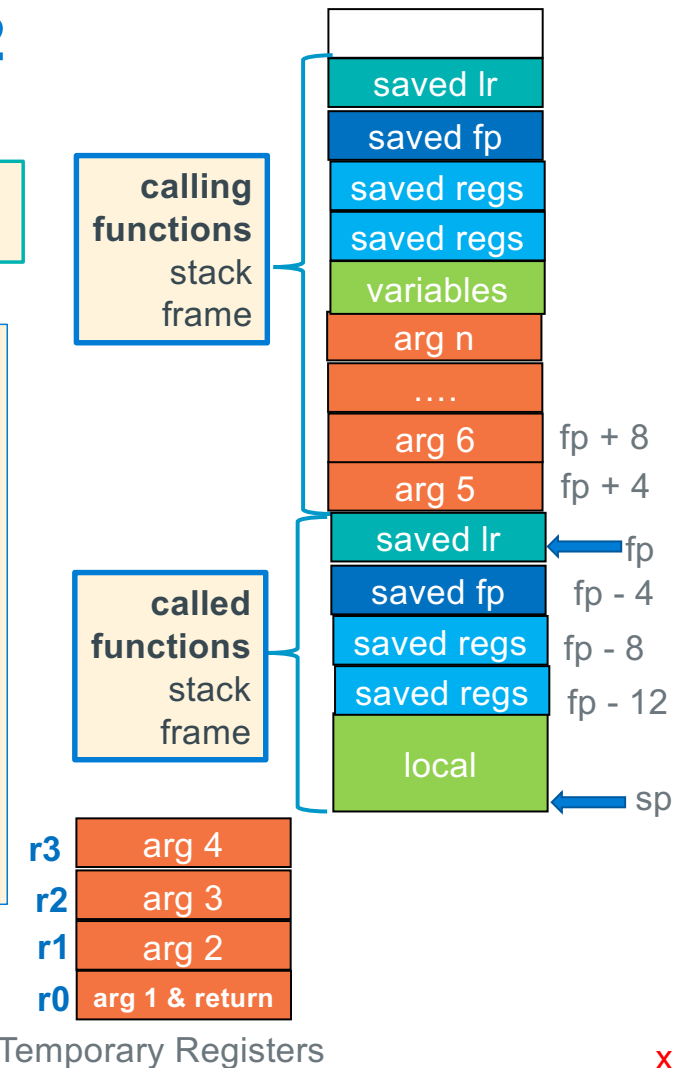


Passing More Than Four Arguments - 2

```
r0 = function(r0, r1, r2, r3, arg5, arg6, ... argn)
      arg1, arg2, arg3, arg4, ...
```

- **Addressing rules**
 - Adding to fp to get arg address in caller's frame
 - Subtracting from fp are addresses in called frame
- Why does it work this way?
- This "algorithm" for finding args was designed to enable languages to have variable argument count functions like:


```
printf("conversion list", arg0, ... argn);
```



Passing More Than Four Arguments – Calling Function

```
r0 = function(r0, r1, r2, r3, arg5, arg6, ... argn)
        arg1, arg2, arg3, arg4, ...
```

- Calling function prior to making the call
 1. Evaluate first four args: place resulting values in r0-r3
 2. Arg 5 and greater are evaluated
 3. Store Arg 5 and greater parameter values on the stack
- **One arg value per slot!** – NO arrays across multiple slots
- chars, shorts and ints are directly stored
- Structs (not always), and arrays are passed via a pointer
- **Pointers** passed as output parameters usually contain an address *that points at* the stack, BSS, data, or heap



Version 1.09

UCSD CSE 30


Aarch32 Assembly – Stack Frames, Bit Ops, Shifts

Week 9

Lecture 27

Keith Muller

Dec Vax 11/780 - 1978



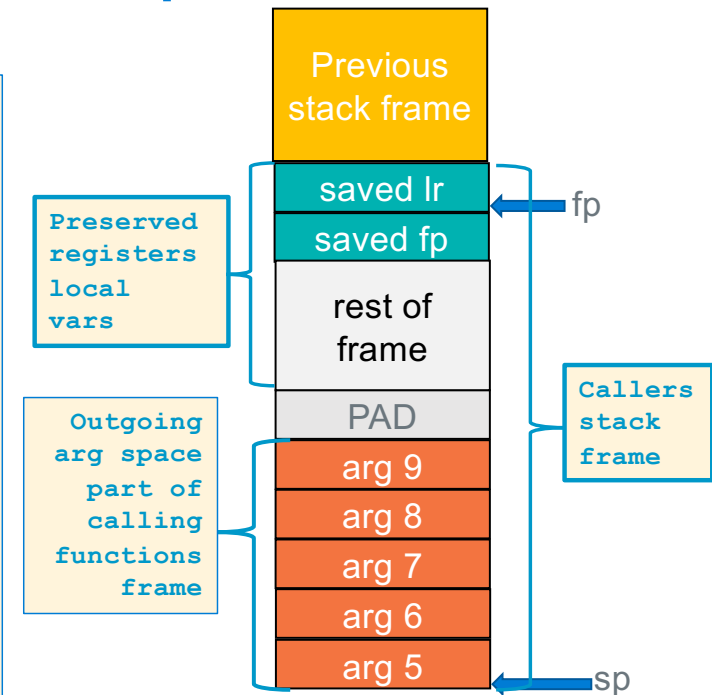
Calling Function: Allocating Stack Parameter Space

At the point of a function call (and obviously at the start of the called function):

1. sp must point at arg5
2. arg5 must be at an 8-byte boundary,
 - a) padding to force arg5 alignment is placed above the last argument the called function is expecting

Approach: Extend the stack frame to include enough space for stack arguments function with the greatest arg count

1. Examine every function call in the body of a function
2. Find the function call with greatest arg count, Determines space needed for outgoing args
3. Add the space needed to the frame layout



Rules: At point of call

1. arg5 must be pointed at by sp
2. SP must be 8-byte aligned

Determining the Passed Parameter Area on The Stack

- Find the function called by main with the largest number of parameters
- That function determines the size of the Passed Parameter allocation on the stack

```
int main(void)
{
    /* code not shown */
    a(g, h);

    /* code not shown */
    sixsum(a1, a2, a3, a4, a5, a6);

    /* code not shown */

    b(q, w, e, r);
    /* code not shown */
}
```

largest arg count is 6
allocate space for $6 - 4 = 2$ arg slots

Passing More than Four Args – Six Arg Example

- Problem: Write and call a function that receives six integers and returns the sum
- First 4 parameters are in register r0 - r3 and the remaining argument are on the stack
- For this example, we will put all the locals on the stack

```
int main(void)
{
    int cnt = sixsum(1, 2, 3, 4, 5, 6);

    printf("the sum is %d\n", cnt);
    return EXIT_SUCCESS;
}
```

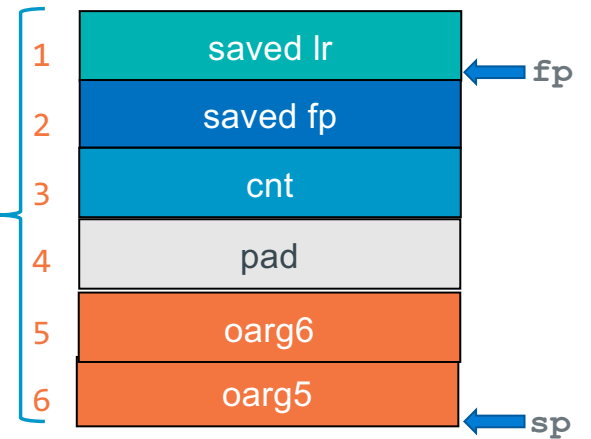
```
int
sixsum(int a1, int a2, int a3, int a4, int a5, int a6)
{
    return a1 + a2 + a3 + a4 + a5 + a6;
}
```

Calling Function > 4 Args - 1

```
int cnt = sixsum(1, 2, 3, 4, 5, 6);
```

```
.equ  FP_OFF,      4  // local base
      // NAME,      SIZE + prev_name
.equ  CNT,          4 + FP_OFF
.equ  PAD,          4 + CNT
.equ  OARG6,        4 + PAD
.equ  OARG5,        4 + OARG6
.equ  FRAMESZ       OARG5 - FP_OFF
```

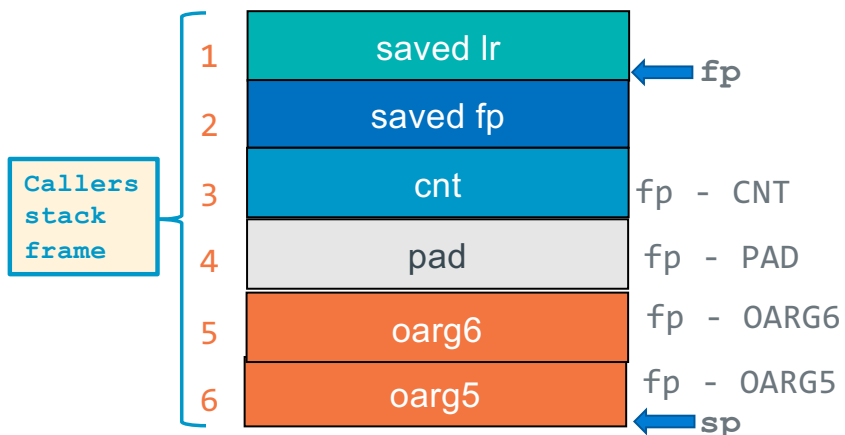
Callers
stack
frame



Calling Function > 4 Args - 2

```
int cnt = sixsum(1, 2, 3, 4, 5, 6);
```

```
.equ  FP_OFF,      4
.equ  CNT,         4 + FP_OFF
.equ  PAD,         4 + CNT
.equ  OARG6,       4 + PAD
.equ  OARG5,       4 + OARG6
.equ  FRAMESZ      OARG5 - FP_OFF
```



```
main:
    push    {fp, lr}
    add     fp, sp, FP_OFF
    sub     sp, sp, FRAMESZ

    mov     r0, 6
    str     r0, [fp, -OARG6]
    mov     r0, 5
    str     r0, [fp, -OARG5]
    mov     r3, 4
    mov     r2, 3
    mov     r1, 2
    mov     r0, 1
    bl      sixsum
    str     r0, [fp, -CNT]
    mov     r1, r0
    ldr     r0, =.Lpfstr
    bl      printf

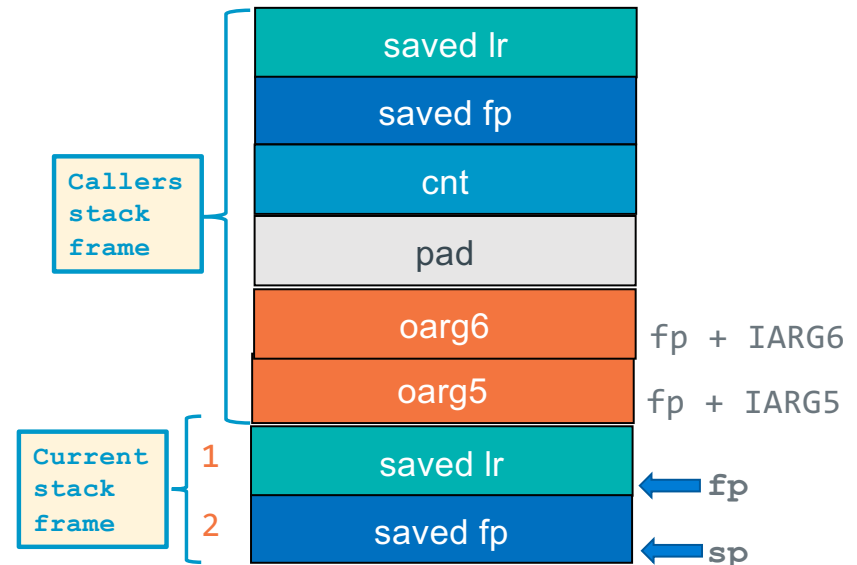
    mov     r0, EXIT_SUCCESS
    sub     sp, fp, FP_OFF
    pop     {fp, lr}
    bx      lr
```

Called Function > 4 Args

```
int sixsum(int a1, int a2, int a3, int a4, int a5, int a6)
    return a1 + a2 + a3 + a4 + a5 + a6;
```

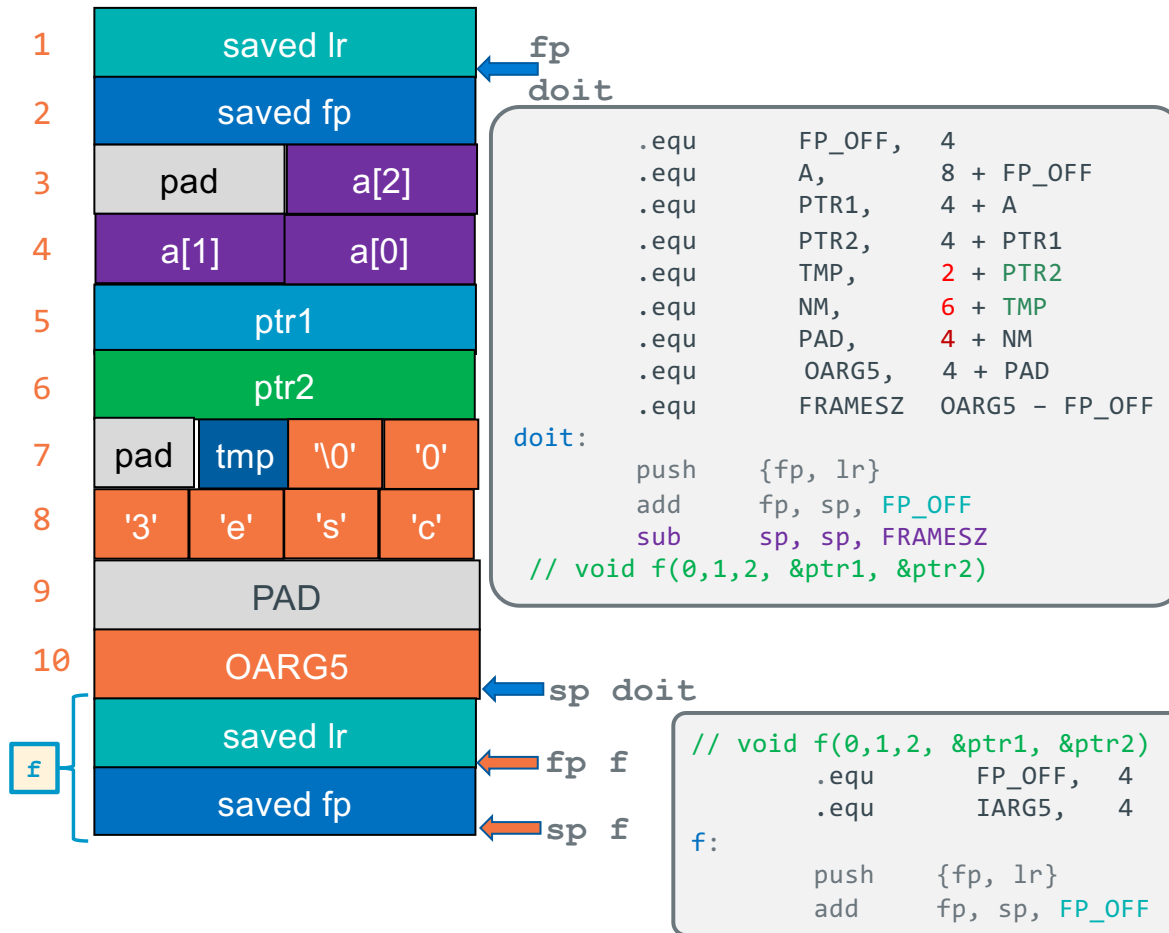
```
.equ  IARG6,      8 // offset into caller's frame
.equ  IARG5,      4 // offset into caller's frame
.equ  FP_OFF,     4 // local base
```

```
sixsum:
    push    {fp, lr}
    add     fp, sp, FP_OFF
    add     r0, r0, r1
    add     r0, r0, r2
    add     r0, r0, r3
    ldr     r1, [fp, IARG5]
    add     r0, r0, r1
    ldr     r1, [fp, IARG6]
    add     r0, r0, r1
    sub     sp, fp, FP_OFF
    pop     {fp, lr}
    bx     lr
```



Recap: Passing pointers

Hint: Useful for PA5



	How to pass
&ptr1	add r3, fp, -PTR1
&ptr2	add r4, fp, -PTR2 str r4, [fp, -OARG5]

Assume that while running, `f()` obtained two pointers from `malloc()` that it returns to `doit` in output `arg4` and `arg5`

	How to change
&ptr1	value to output is in r6 str r6, [r3]
&ptr2	value to output is in r7 ldr r8, [fp, IARG5] str r7, [r8]

Bitwise (Bit to Bit) Operators in C

output = ~a;

a	~a
0	1
1	0

output = a & b;

a	b	a & b
0	0	0
0	1	0
1	0	0
1	1	1

& with 1 to let a bit through
& with 0 to set a bit to 0

output = a | b;

a	b	a b
0	0	0
0	1	1
1	0	1
1	1	1

| with 1 to set a bit to 1
| with 0 to let a bit through

output = a ^ b; //EOR

a	b	a ^ b
0	0	0
0	1	1
1	0	1
1	1	0

^ with 1 will flip the bit
^ with 0 to let a bit through

Bitwise
NOT

~	1100

	0011

Bitwise
AND

	0110
&	1100

	0100

Bitwise
OR

	0110
	1100

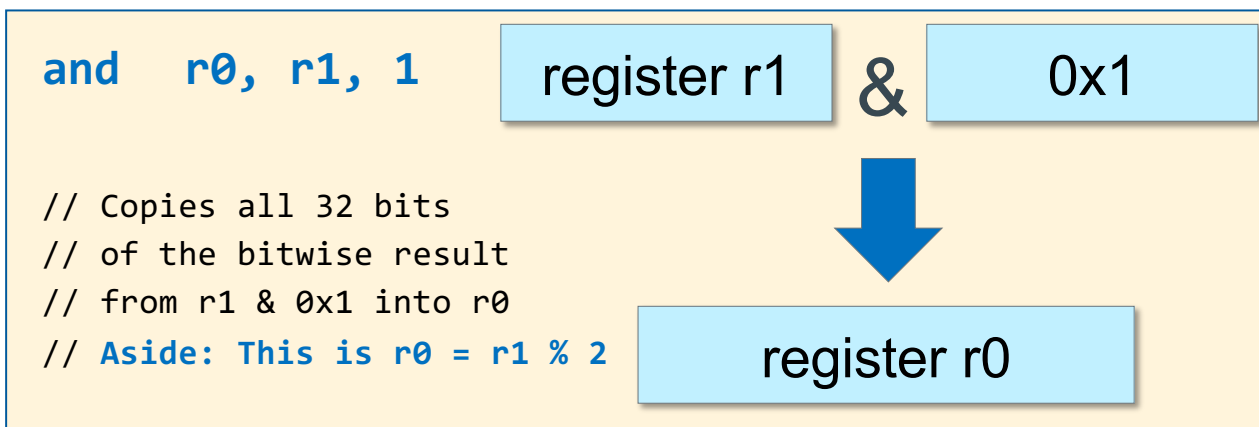
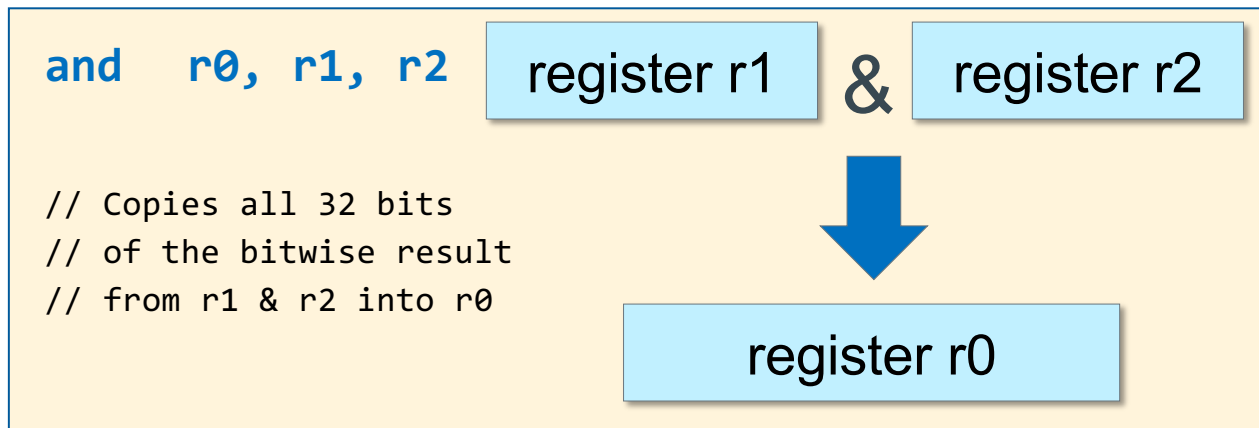
	1110

Bitwise
EOR



	0110
^	1100

	1010

First Look: **AND** Registers



Bitwise Instructions



`<op> Rd, Rn, constant // Rd = Rn <op> constant`
`<op> Rd, constant // Rd = Rd <op> constant`
`<op> Rd, Rn, Rm // Rd = Rn <op> Rm`

Bytes: $0 \leq \text{imm8} \leq 255$ + values from "rotating" rot 4 bits

Bitwise <code><op></code> description	<code><op></code> Syntax	Operation
Bitwise AND	<code>and Rd, Rn, Op2</code>	$R_d \leftarrow R_n \& Op2$
Bit Clear each bit in Op2 that is a 1, the same bit in R_d , is cleared	<code>bic Rd, Rn, Op2</code>	$R_d \leftarrow R_n \& \sim Op2$
Bitwise OR	<code>orr Rd, Rn, Op2</code>	$R_d \leftarrow R_n Op2$
Exclusive OR	<code>eor Rd, Rn, Op2</code>	$R_d \leftarrow R_n \wedge Op2$

Bit Masks: Masking - 1

- Bit masks access/modify specific bits in memory
- Masking act of applying a mask to a value
- **or**: 0 passes bit unchanged, 1 sets bit to 1
- **eor**: 0 passes bit unchanged, 1 inverts the bit
- **bic**: 0 passes bit unchanged, 1 clears it
- **and**: 0 clears the bit, 1 passes bit unchanged

mask force lower 16 bits to 1 "**mask on**" operation

```
orr  r1, r2, r3
```

DATA: r2 0xab ab ab 77

MASK: r3 0x00 00 ff ff lower half to 1

RSLT: r1 0xab ab ff ff

mask to invert the lower 8-bits "**bit toggle**" operation

```
eor  r1, r2, r3
```

DATA: r2 0xab ab ab 77

MASK: r3 0x00 00 00 ff flip LSB bits

RSLT: r1 0xab ab ab 88

MASK: r3 0x00 00 00 ff apply a 2nd time

RSLT: r1 0xab ab ab 77 original value!

x

Bit Masks: Masking - 2

mask to **extract top 8 bits** of r2 into r1

and r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0xff 00 00 00

RSLT: r1 0xab 00 00 00

mask to query the status of a bit **"bit status"** operation

and r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0x00 00 00 01 is bit 0 set?

RSLT: r1 0x00 00 00 01 (0 if not set)

mask to force lower 8 bits to 0 **"mask off"** operation

and r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0xff ff ff 00 clear LSB

RSLT: r1 0xab ab ab 00

clear bit 5 to a 0 without changing the other bits

bic r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0x00 00 00 20 clear bit 5 (0010)

RSLT: r1 0xab ab ab 57

Bit Masks: Masking - 3

mask to get **1's complement** operation
(like mvn)

eor r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0xff ff ff ff

RSLT: r1 0x54 54 54 88

remainder (mod): num \% d where $n \geq 0$ and $d = 2^k$

$\text{mask} = 2^k - 1$ so for mod 2, $\text{mask} = 2 - 1 = 1$

and r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0x00 00 00 01 (mod 2 even or odd)

RSLT: r1 0x00 00 00 01 (odd)

remainder (mod): num \% d where $n \geq 0$ and $d = 2^k$

$\text{mask} = 2^k - 1$ so for mod 16, $\text{mask} = 16 - 1 = 15$


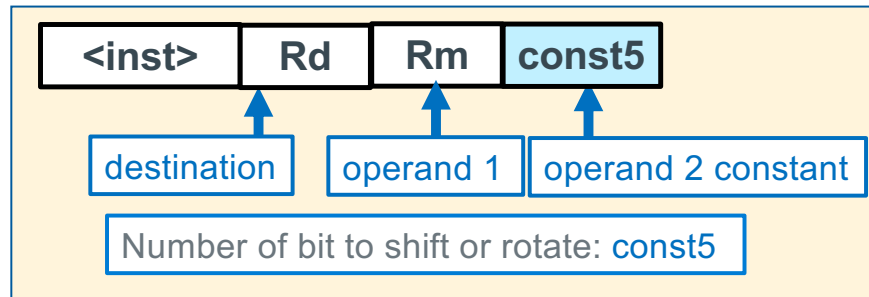
and r1, r2, r3

DATA: r2 0xab ab ab 77

MASK: r3 0x00 00 00 0f (mod 16)

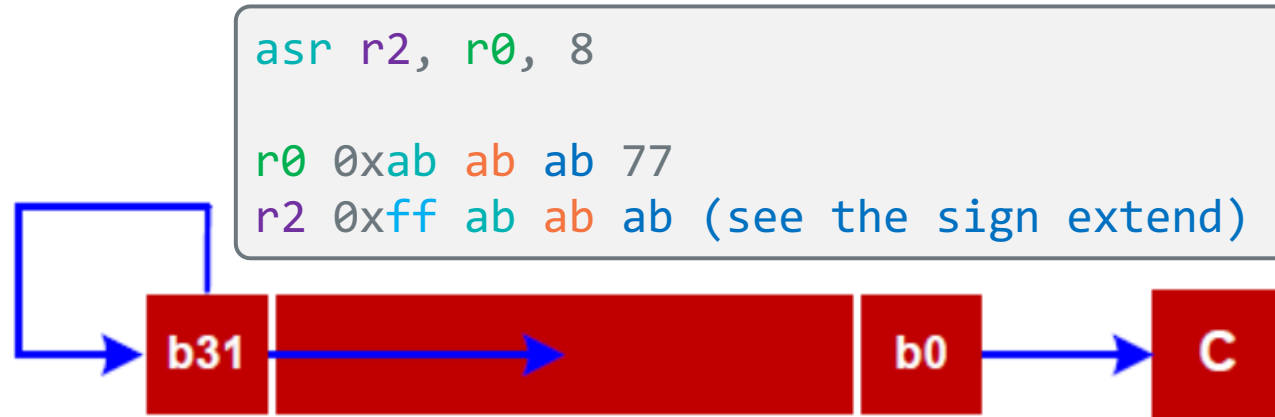
RSLT: r1 0xab 00 00 07 (if 0: divisible by)

Shift and Rotate Instructions



Instruction	Syntax	Operation	Notes	Diagram
Logical Shift Left	LSL $R_d, R_m, const5$	$R_d \leftarrow R_m \ll const5$	Zero fills shift: 0 - 31	
	LSL R_d, R_m, R_s	$R_d \leftarrow R_m \ll R_s$		
Logical Shift Right	LSR $R_d, R_m, const5$	$R_d \leftarrow R_m \gg const5$	Zero fills shift: 1 - 32	
	LSR R_d, R_m, R_s	$R_d \leftarrow R_m \gg R_s$		
Arithmetic Shift Right	ASR $R_d, R_m, const5$	$R_d \leftarrow R_m \ggg const5$	Sign extends shift: 1 - 32	
	ASR R_d, R_m, R_s	$R_d \leftarrow R_m \ggg R_s$		
Rotate Right	ROR $R_d, R_m, const5$	$R_d \leftarrow R_m \text{ ror } const5$	right rotate rot: 0 - 31	
	ROR R_d, R_m, R_s	$R_d \leftarrow R_m \text{ ror } R_s$		

Shift & Rotate Operations



Test for sign
-1 if r0 negative

```
asr r2, r0, 31  
  
r0 0xab ab ab 77  
r2 0xff ff ff ff
```

Test for sign
0 if r0 positive

```
asr r2, r0, 31  
  
r0 0x7b ab ab 77  
r2 0x00 00 00 00
```


Shift & Rotate Operations




```
lsr r2, r0, 8
```

```
r0 0xab ab ab 77
r2 0x00 ab ab ab
```



```
lsl r2, r0, 8
```

```
r0 0xab ab ab 77
r2 0xab ab 77 00
```




```
ror r2, r0, 8
```

```
r0 0xab ab ab 77
r2 0x77 ab ab ab
```


Extracting Unsigned Bitfields

- Move byte 2 in r0 to byte 0 in r1

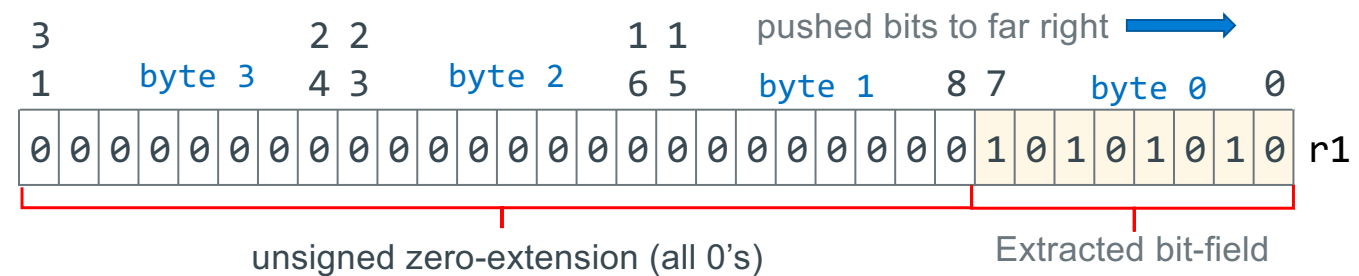
Hint: Useful for PA5



lsl r1, r0, 8




lsr r1, r1, 24




Extracting Signed Bitfields

- Move byte 2 in r0 to byte 0 in r1




next shift left = 8

lsl r1, r0, 8



next shift right = 24

asr r1, r1, 24



signed extend (all 1's)

Extracted
bit-field

Inserting Bitfields – Inserting Source Field into Destination Field

Task: Insert source into destination

a	b	a b
0	0	0
0	1	1
1	0	1
1	1	1

Approach


(1) isolate source field

(2) clear destination field

(3) Bitwise **or** together

orr r1, r1, r2

results in




Creating a Mask

option #1 (1 mask)

```
ldr    r3, =0x00ffff0000
```

for a 0 mask

```
ldr    r3, =0xff0000ffff
```



```
option #2 (1 mask)
```


small mask

```
mov    r3, 255
```

```
lsl    r3, r3, 16
```

or do

```
ror    r3, r3, 16
```



next shift left = 16 bits



Creating a Mask- 0 mask


option #3
any size field

```
mov    r3, -1
```


number of bits you need in
the mask, 8 for example

```
asr    r3, r3, 8
```

```
ror    r3, r3, 8
```



next shift right = 8 bits



next rotate right = 8 bits



Creating a Mask- 1 mask

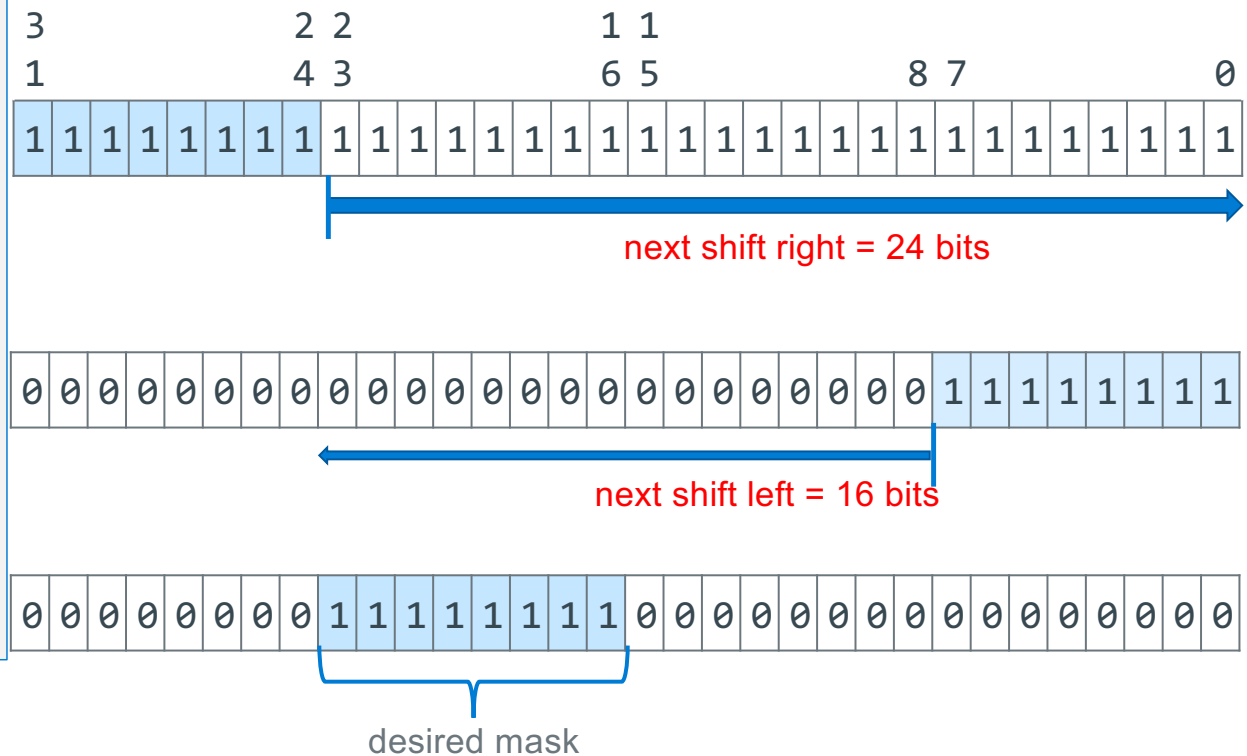
option #3
any size field

```
mov    r3, -1
```


32 - number of bits you
need in the mask, 8 for
example is mask size

```
lsr    r3, r3, 24
```

```
lsl    r3, r3, 16
```




Inserting Bitfields – Isolating the Source Field




isolate source field

`lsl r2, r0, 24`

`lsr r2, r2, 8`



Inserting Bitfields – Clearing the Destination Field




create a 1 mask




clear the
destination field

bic r1, r1, r3




Inserting Bitfields – Combining Isolated Source and Cleared Destination

isolated source



field cleared in
destination



inserted field
orr r1, r1, r0




Masking Summary


Isolate a field: Use **and** with a **mask** of one's surrounded by zero's to select the bits that have a 1 in the mask, all other bits will be set to zero



Clear a field: Use **and** with a mask of zero's surrounded by one's to select the bits that have a 1 in the mask, all other bits will be set to zero



Isolate a field: Use **lsl** and **lsr** to get a field surrounded by zeros



Reference For PA5: C Stream Functions Opening Files

```
FILE *fopen(char filename[], const char mode[]);
```

- Opens a stream to the specified file in specified file access mode
 - returns NULL on failure – **always check the return value; make sure the open succeeded!**

- Mode is a string that describes the actions that can be performed on the stream:

"r" Open for reading.

The stream is positioned at the beginning of the file. Fail if the file does not exist.

"w" Open for writing.

The stream is positioned at the beginning of the file. Create the file if it does not exist.

"a" Open for writing.

The stream is positioned at the end of the file. Create the file if it does not exist.

Subsequent writes to the file will always be at current end of file.

- An optional "+" following "r", "w", or "a" opens the file for both reading and writing

Reference: C Stream Functions Closing Files and Usage

```
int fclose(FILE *stream) ;
```

- Closes the specified stream, forcing output to complete (eventually)
 - returns EOF on failure (often ignored as no easy recovery other than a message)
- Usage template for `fopen()` and `fclose()`
 1. Open a file with `fopen()` **always** checking the return value
 2. do i/o – keep calling stdio io routines
 3. close the file with `fclose()` when done with that I/O stream

C Stream Functions Array/block read/write

- These do not process contents they simply **transfer** a fixed number of bytes to and from a buffer passed to them
- `size_t fwrite(void *ptr, size_t size, size_t count, FILE *stream);`
 - Writes an array of *count elements* of *size* bytes from *stream*
 - *Updates the write file pointer forward by the number of bytes written*
 - returns number of elements written
 - error is short element count or 0
- `size_t fread(void *ptr, size_t size, size_t count, FILE *stream);`
 - Reads an array of *count elements* of *size* bytes from *stream*
 - *Updates the read file pointer forward by the number of bytes read*
 - returns number of elements read, **EOF is a return of 0**
 - error is short element count or 0
- **I almost always set size to 1 to return bytes read/written**

C fread/fwrite Example - 1

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#define BFSZ      8192 /* size of read */
int main(void)
{
    char fbuf[BFSZ];
    FILE *fin, *fout;
    size_t readlen;
    size_t bytes_copied = 0;
    retval = EXIT_SUCCESS;

    if (argc != 3){
        fprintf(stderr, "%s requires two args\n", argv[0]);
        return EXIT_FAILURE;
    }
    /* Open the input file for read */
    if ((fin = fopen(argv[1], "r")) == NULL) {
        fprintf(stderr, "fopen for read failed\n");
        return EXIT_FAILURE;
    }
    /* Open the output file for write */
    if ((fout = fopen(argv[2], "w") == NULL) {
        fprintf(stderr, "fopen for write failed\n");
        fclose(fin);
        return EXIT_FAILURE;
    }
}
```

To handle
bytes moved

```
% ls -ls ZZZ
ls: ZZZ: No such file or directory
% ./a.out cp.c ZZZ
bytes copied: 1122
% ls -ls cp.c ZZZ
8 -rw-r--r--  1 kmuller  staff  1122 Jul  2 08:51 ZZZ
8 -rw-r--r--  1 kmuller  staff  1122 Jul  2 08:49 cp.c
```


C fread/fwrite Example - 2

```
/* Read from the file, write to fout */  
  
while ((readlen = fread(fbuf, 1, BUFSIZ, fin)) > 0) {  
    if (fwrite(fbuf, 1, readlen, fout) != readlen) {  
        fprintf(stderr, "write failed\n");  
        retval = EXIT_FAILURE;  
        break;  
    }  
    bytes_copied += readlen; //running sum bytes copied  
}  
  
if (retval == EXIT_FAILURE)  
    printf("Failure Copy did not complete only ");  
printf("Bytes copied: %zu\n", bytes_copied);  
  
fclose(fin);  
fclose(fout);  
  
return retval;  
}
```

By using an element size of 1 with a char buffer, this is byte I/O

Capture the bytes read so you know how many bytes to write

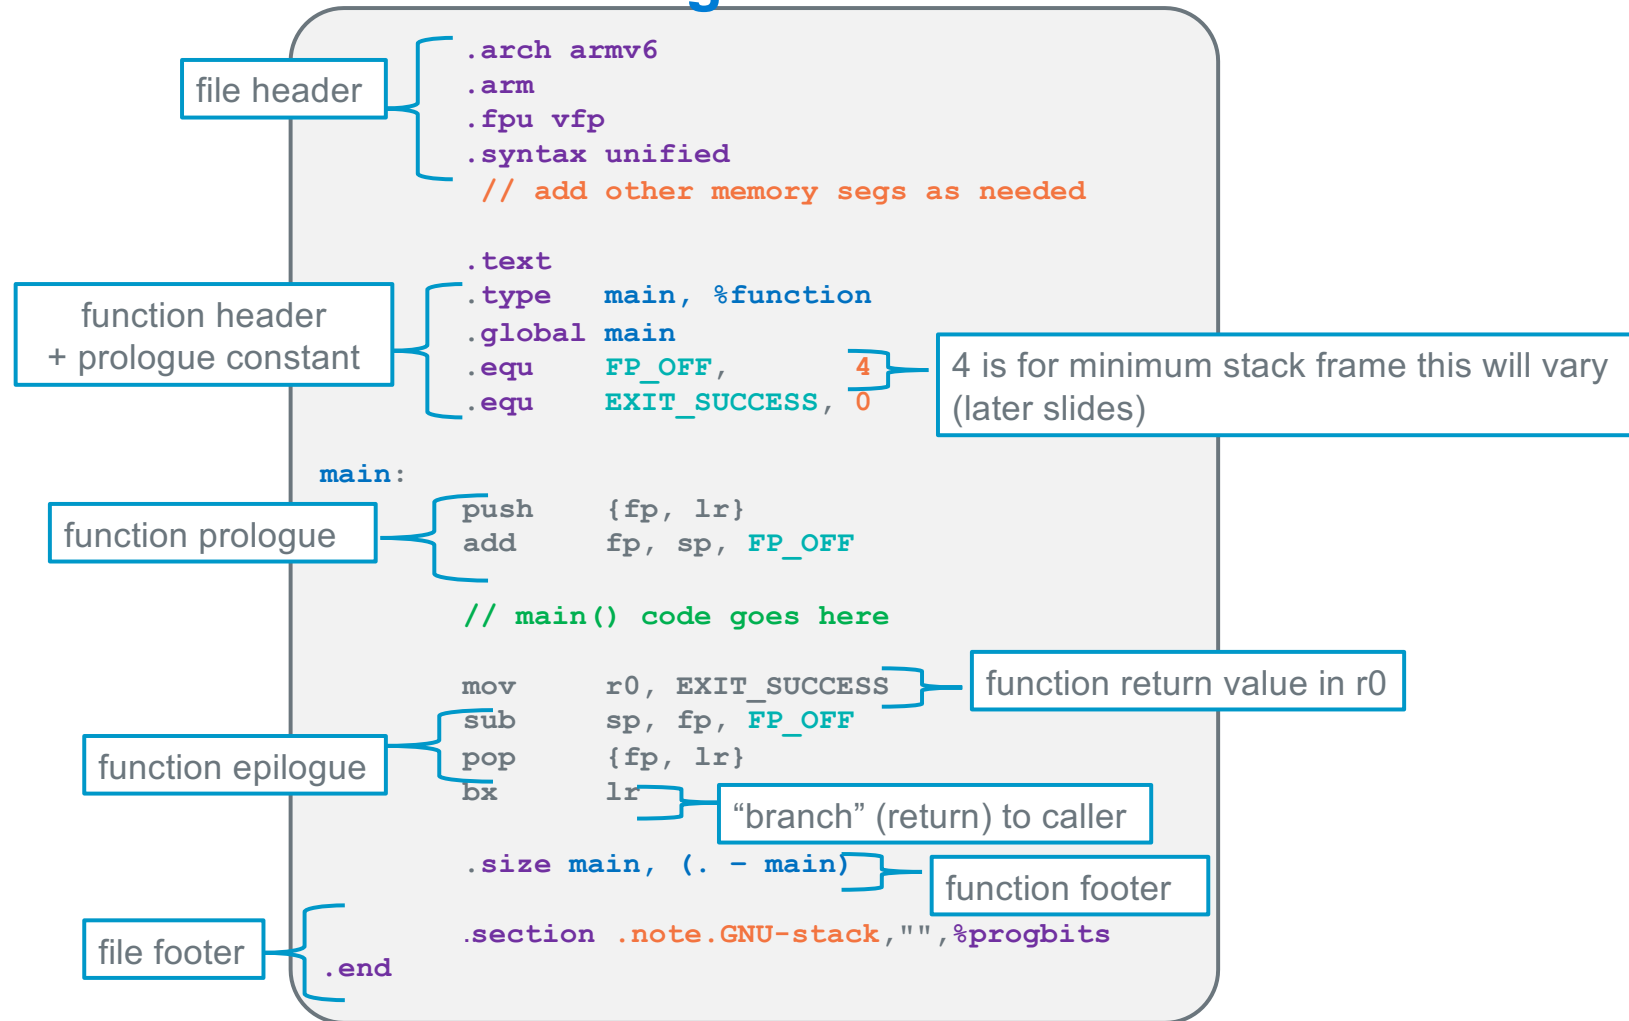
unless file length is an exact multiple of BUFSIZ, the last fread() will always be less than BUFSIZ which is why you write readln



Jargon: the last record is often called the "runt"

Extras

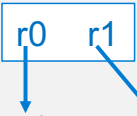
main.S Source File Showing a minimum stack frame



putchar/getchar Setting up and Usage

```
#include <stdio.h>
#include <stdlib.h>
int
main(void)
{
    int c;
    int count = 0;

    while ((c = getchar()) != EOF) {
        putchar(c);
        count++;
    }
    printf("Echo count: %d\n", count);
    return EXIT_SUCCESS;
}
```



```
.extern getchar
.extern putchar
.section .rodata
.Lfstr: .string "Echo count: %d\n"
.text
.equ    EOF,      -1
.type   main, %function
.global main
.equ    FP_OFF,    12
.equ    EXIT_SUCCESS, 0
main:   push      {r4, r5, fp, lr}
        add       fp, sp, FP_OFF
        mov       r4, 0 //r4 = count

/* while loop code will go here */
.Ldone:
        mov       r1, r4 // count
        ldr       r0, =.Lfstr
        bl        printf
        mov       r0, EXIT_SUCCESS
        sub       sp, fp, FP_OFF
        pop       {r4, r5, fp, lr}
        bx        lr
        .size main, (. - main)
```

Putchar/getchar: The while loop

```
#include <stdio.h>
#include <stdlib.h>
int
main(void)
{
    int c;
    int count = 0;

    while ((c = getchar()) != EOF) {
        putchar(c);
        count++;
    }
    printf("Echo count: %d\n", count);
    return EXIT_SUCCESS;
}
```

initialize count

pre loop test with a call to getchar()
if it returns EOF in r0 we are done

echo the character read with getchar and
then read another and increment count

did getchar() return EOF if not loop

saw EOF, print count

```
mov    r4, 0    //count
bl     getchar
cmp    r0, EOF
beq    .Ldone

.Lloop:
bl     putchar
bl     getchar
add    r4, r4, 1
cmp    r0, EOF
bne    .Lloop

.Ldone:
mov    r1, r4
ldr    r0, =pfstr
bl     printf
```

File header and footers are not shown

printing error messages in assembly

```
.Lmsg0: .string "Read failed\n"
        ldr    r0, =.Lmsg0           // read failed print error
        bl     errmsg
```

```
        // int errmsg(char *errmsg)
        // writes error messages to stderr
        .type  errmsg, %function      // define to be a function
        .equ   FP_OFF, 4              // fp offset in stack frame
errmsg:
        push   {fp, lr}               // stack frame register save
        add    fp, sp, FP_OFF         // set the frame pointer

        mov    r1, r0
        ldr     r0, =stderr
        ldr     r0, [r0]
        bl     fprintf
        mov     r0, EXIT_FAILURE      // Set return value
        sub     sp, fp, FP_OFF        // restore stack frame top
        pop     {fp, lr}              // remove frame and restore
        bx      lr                    // return to caller
        // function footer
        .size   errmsg, (. - errmsg)  // set size for function
```

Reference: Registers and Flags – Programmers View

Parameters/Return/Scratch Registers

Register Content is **NOT protected** across function calls

Function Arguments and return values

`r0 = function(r0, r1, r2, r3)`

for very special cases:


(64 bits) `r1, r0 = function(r0, r1, r2, r3)`



Preserved registers

Register Content **is protected** across function calls

Used for passing system call #'s to the OS



Hardware & Restricted Use Registers

Frame Pointer `r11/fp`

Intra Procedure Call Register `r12/ip`

Linker uses to make "long" function calls

Stack Pointer `r13/sp`

Link Register `r14/lr`

Program Counter – **Do not use** `r15/pc`

CSPR Register Flags

N Negative Flag

Z Zero Flag

C Carry Flag

V Overflow Flag

Parameter Registers



Preserved Registers



Scratch Registers



Special Use Registers



Bitwise versus C Boolean Operators

Meaning	Operator	Operator	Meaning
Boolean AND	a && b	a & b	Bitwise AND
Boolean OR	a b	a b	Bitwise OR
Boolean NOT	!b	~b	Biwise NOT

Boolean operators **act on the entire value not the individual bits**

& versus &&

`0x10 & 0x01 = 0x00 (bitwise)`

`0x10 && 0x01 = 0x01 (Boolean)`

! versus ~

`~0x01 = 0xfffffffffe (bitwise)`

`!0x01 = 0x0 (Boolean)`